

VE281

Data Structures and Algorithms

Priority Queues and Heaps

Learning Objectives:

- Know what a priority queue is
- Know what a min heap is
- Know how a min heap performs enqueue and extractMin operations
- Know how to efficiently initialize a min heap

Outline

- ❖ Priority Queue
- ❖ Min Heap and Its Operations
- ❖ Min Heap Initialization and Application

Priority Queues

- ❖ Two kinds of priority queues:
 - ❖ Min priority queue.
 - ❖ Max priority queue.
- ❖ We will focus on **min priority queue**.
 - ❖ The max priority queue is similar.

What Is Min Priority Queue?

- ❖ A collection of items.
- ❖ Each item has a key (or “**priority**”).
- ❖ Support the following operations:
 - ❖ **isEmpty**
 - ❖ **size**
 - ❖ **enqueue**: put an item into the priority queue.
 - ❖ **dequeueMin**: remove element with **min** key.
 - ❖ **getMin**: get item with **min** key.

Applications of Priority Queue

- ❖ Banking services
 - ❖ VIP customer who arrives later gets served first.
- ❖ Network bandwidth management
 - ❖ The prioritized traffic, such as real-time data, is forwarded with the least delay once it reaches the network router.
- ❖ Discrete event simulation
 - ❖ One event happening triggers a few others, which are put into a queue.
 - ❖ Simulating in the order of the **beginning time** of the events.

Min Priority Queue: Implementation

- ❖ A collection of items.
- ❖ Each item has a key (or “**priority**”).
- ❖ Support the following operations:
 - ❖ **isEmpty**
 - ❖ **size**
 - ❖ **enqueue**: put an item into the priority queue.
 - ❖ **dequeueMin**: remove element with **min** key.
 - ❖ **getMin**: get item with **min** key.

What's the time complexity for an unsorted array-based implementation?

Priority Queue Implemented with Heap

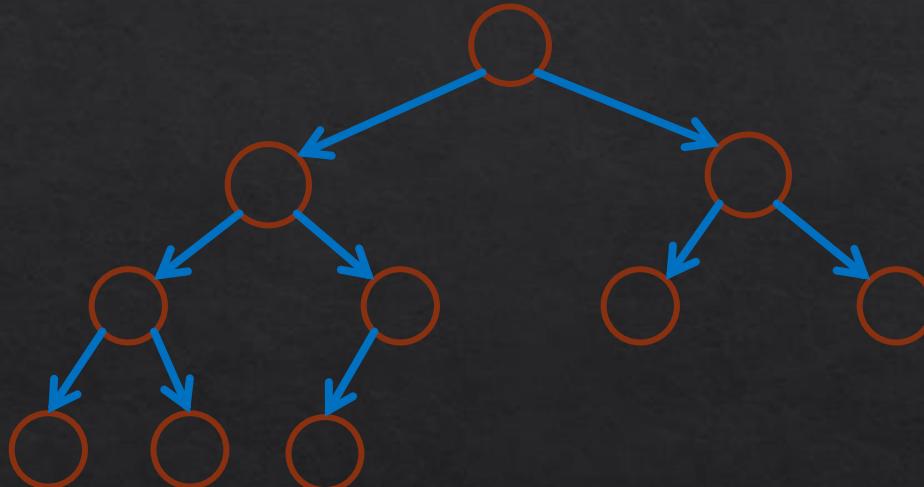
- ❖ Priority queues are most commonly implemented using **Binary Heaps** (will be shown soon).
- ❖ Complexity of the operation using heap implementation:
 - ❖ **isEmpty**, **size**, and **getMin** are $O(1)$ time complexity in the worst case.
 - ❖ **enqueue** and **dequeueMin** are $O(\log n)$ time complexity in the worst case, where n is the size of the priority queue.

Outline

- ❖ Priority Queue
- ❖ Min Heap and Its Operations
- ❖ Min Heap Initialization and Application

Binary Heap

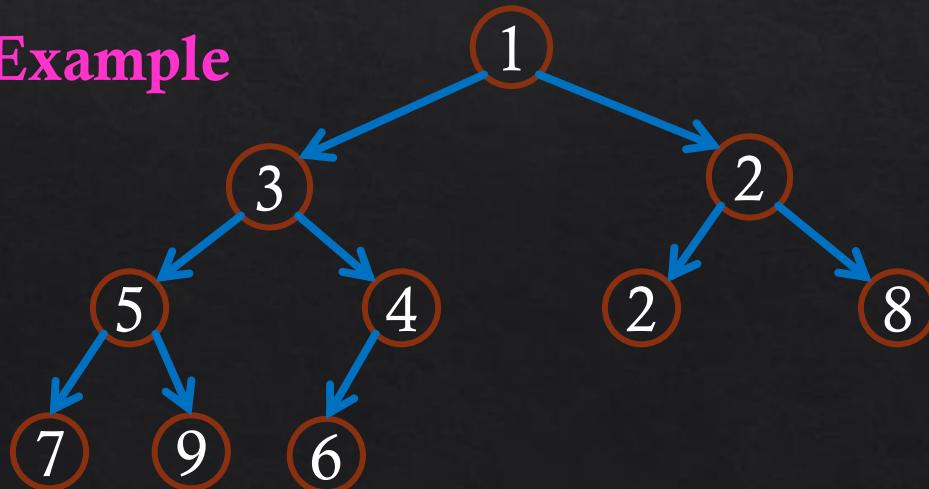
- ❖ A **binary heap** is a **complete binary tree**.



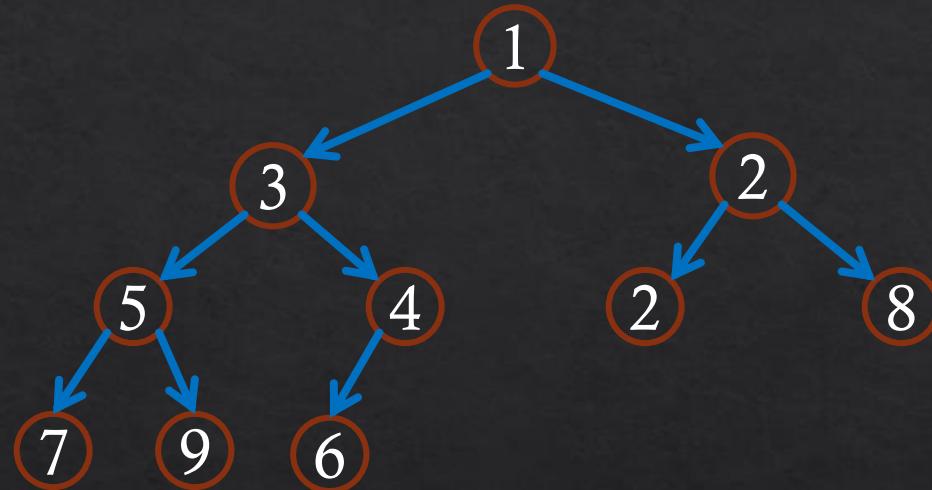
Min Heap

- ❖ A min heap is
 - ❖ a **binary heap**, and
 - ❖ a tree where for **any** node v , the key of v is smaller than or equal to (\leq) the keys of any **descendants** of v .
- ❖ Property: The key of the root of **any** subtree is always the smallest among all the keys in that subtree.

Example



Min Heap



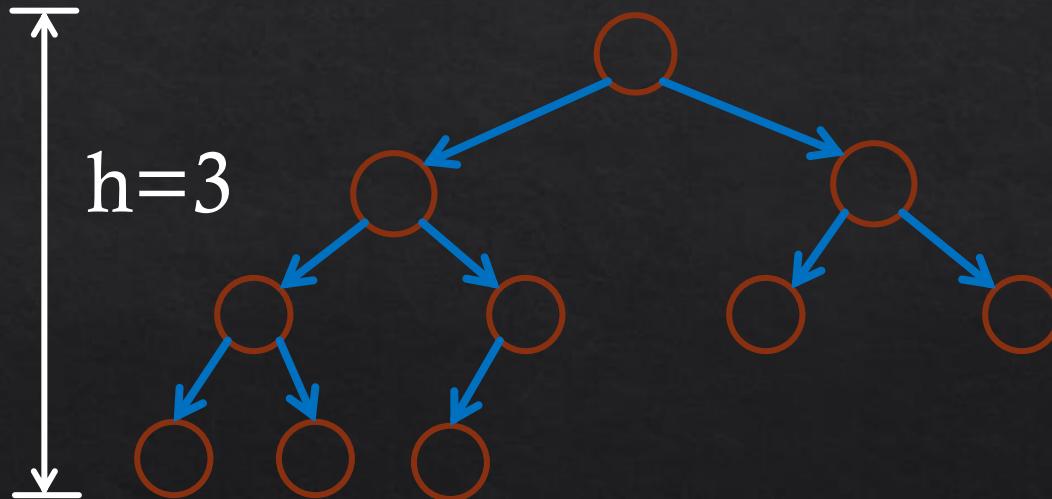
- ❖ However, the keys of nodes **across** subtrees have no required relationship.
- ❖ **Binary heaps** are different from **binary search trees**, which we will show in future lectures.



What's the Height of a Heap of n Nodes?

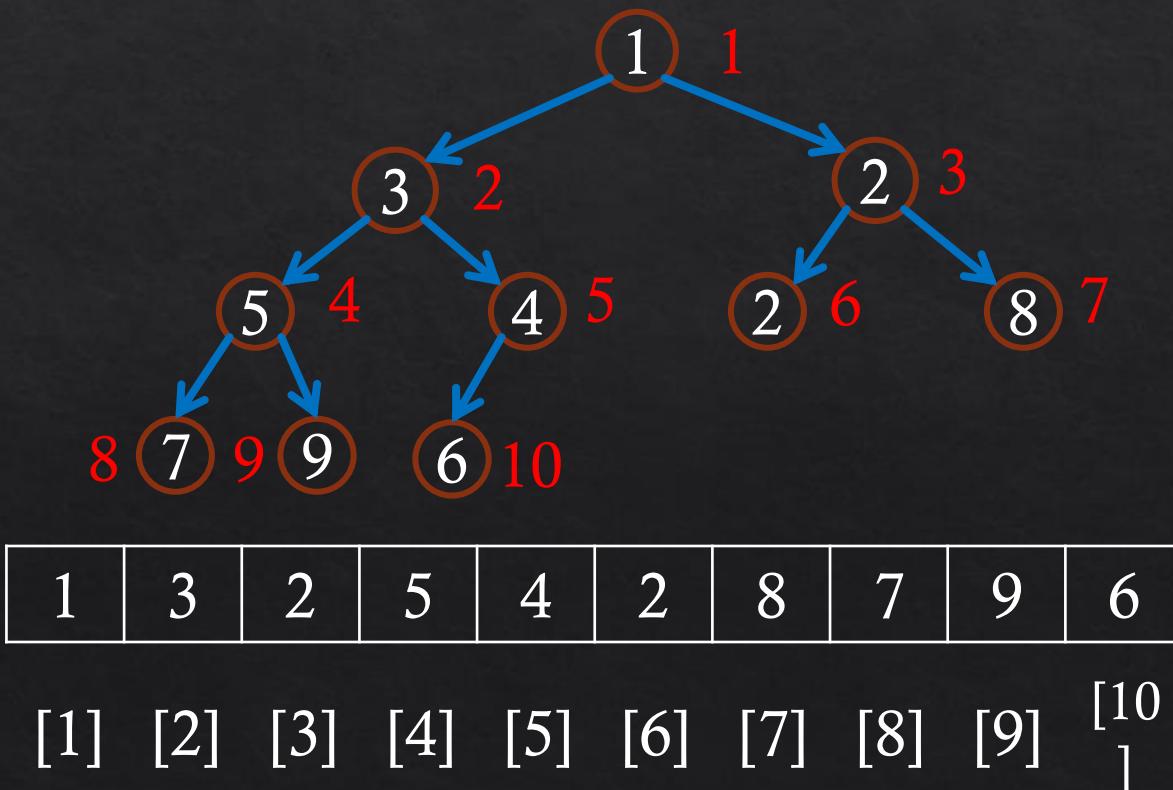
Select **all** the correct answers:

- A. $\lceil \log_2(n + 1) \rceil - 1$
- B. $\lceil \log_2 n \rceil - 1$
- C. $\lceil \log_2(n + 1) \rceil - 1$
- D. $\lceil \log_2 n \rceil$

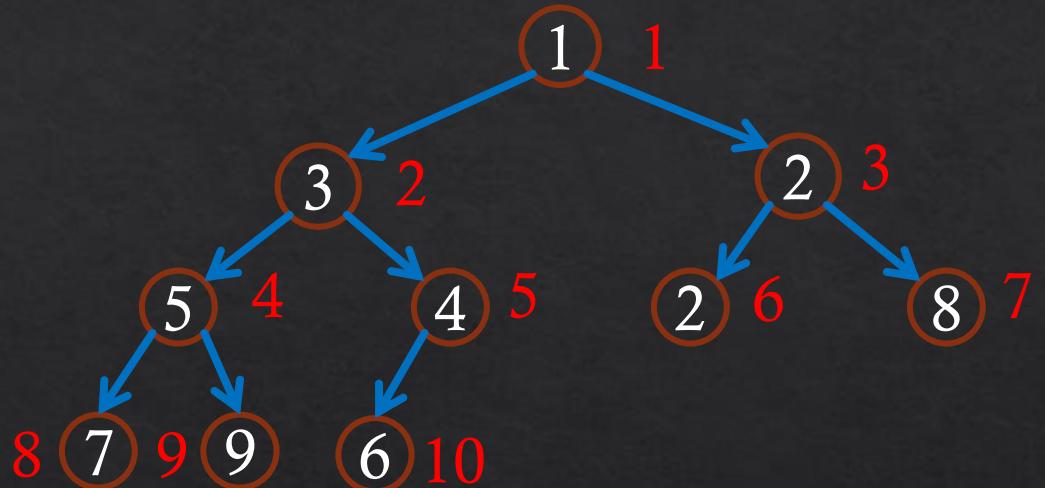


Binary Heap Implementation as an Array

- ❖ Store the elements in an array in the order produced by a level-order traversal.
- ❖ The first element is stored at index 1.



Index Relation



Index relation allows us to move up and down a heap easily.

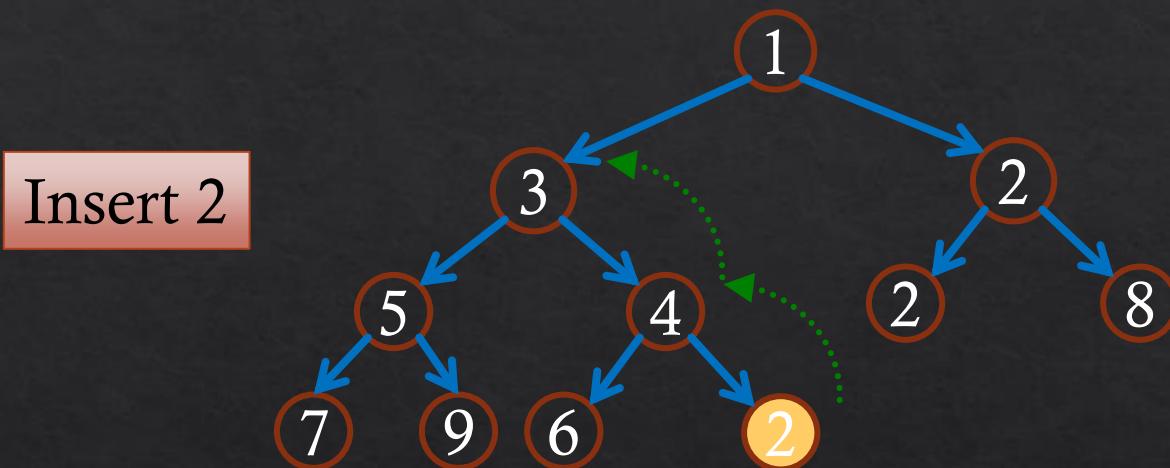
- ◊ A node at index i ($i \neq 1$) has its parent at index $\lfloor i/2 \rfloor$.
- ◊ Assume the number of nodes is n . A node at index i ($2i \leq n$) has its left child at $2i$.
 - ◊ If $2i > n$, it has no left child.
- ◊ A node at index i ($2i + 1 \leq n$) has its right child at $2i + 1$.
 - ◊ If $2i + 1 > n$, it has no right child.

Min Heap Implementation

- ❖ We also have a **size** variable to keep the number of nodes in the heap.
 - ❖ The heap elements are stored in `heap[1]`, `heap[2]`, ..., `heap[size]`.
- ❖ Operations
 - ❖ `isEmpty: return size==0;`
 - ❖ `size: return size;`
 - ❖ `getMin: return heap[1];`

Procedure of enqueue

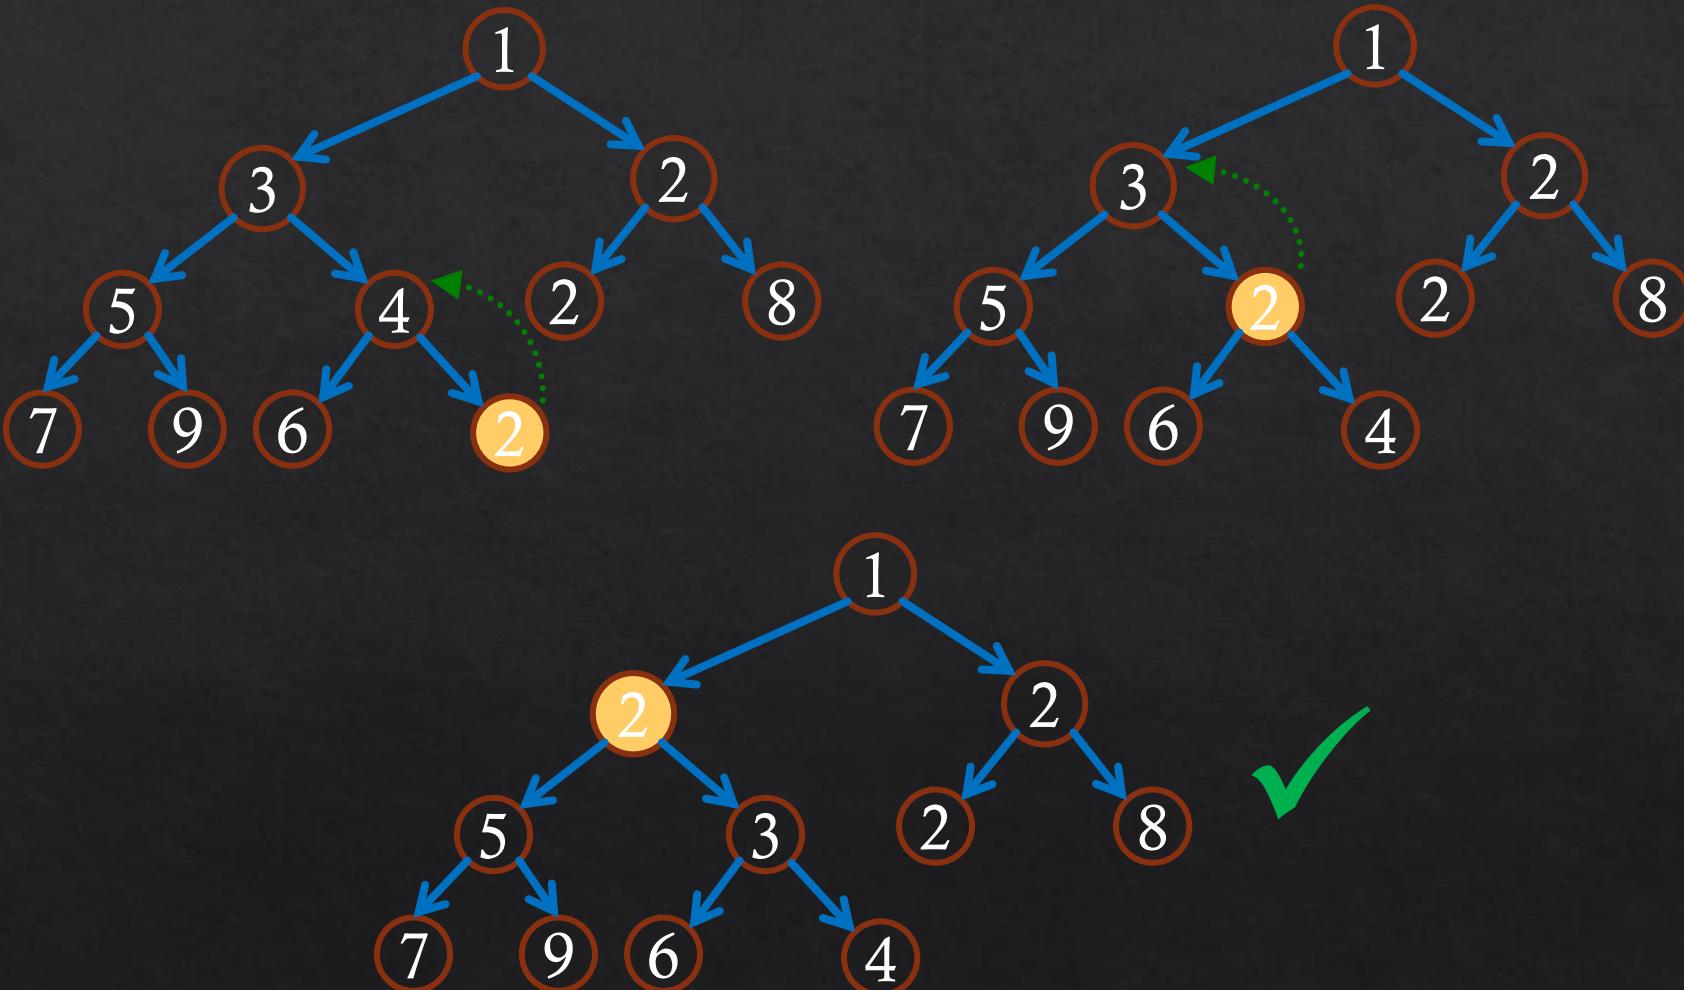
- ◊ Insert **newItem** as the rightmost leaf of the tree.



```
heap[++size] = newItem;
```

- ◊ The tree may no longer be a heap at this point!
- ◊ **Percolate up** **newItem** to an appropriate spot in the heap to restore the heap property.

Percolate Up Illustration



Percolate Up Code

```
void minHeap::percolateUp(int id) {  
    while(id > 1 && heap[id/2] > heap[id]) {  
        swap(heap[id], heap[id/2]);  
        id = id/2;  
    }  
}
```

- ◊ Pass index (**id**) of array element that needs to be percolated up.
- ◊ Swap the given node with its parent and move up to parent until:
 - ◊ we reach the root at position 1, or
 - ◊ the parent has a smaller or equal key.

enqueue

Code

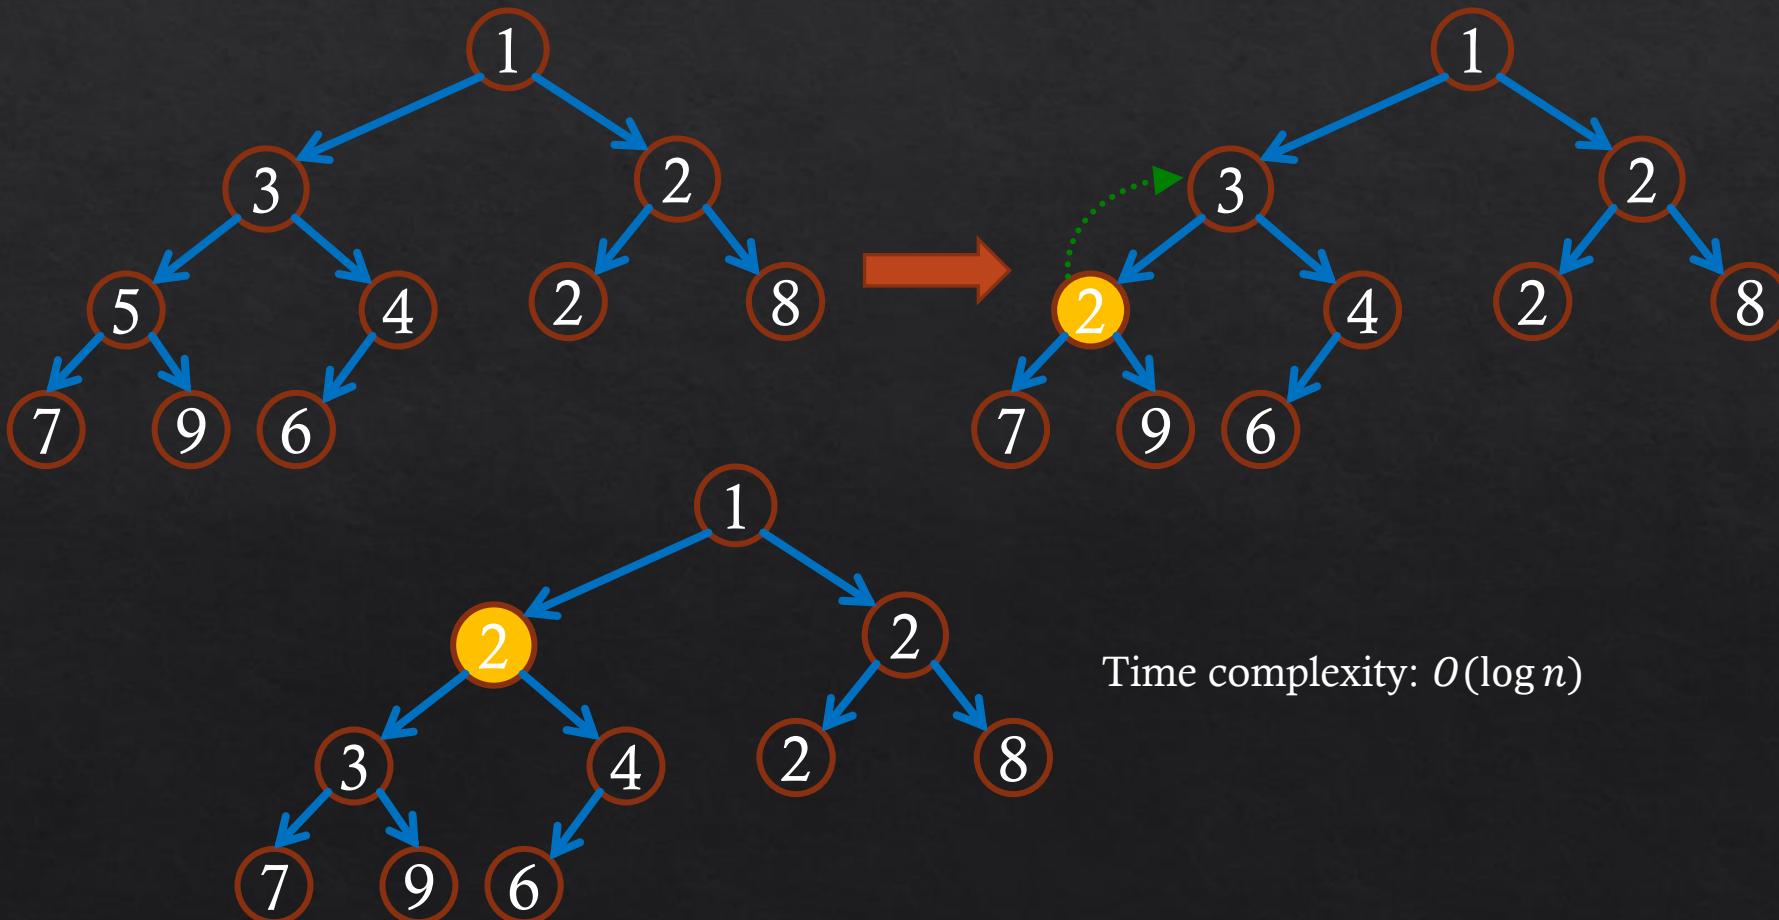
```
void minHeap::enqueue(Item newItem) {  
    heap[++size] = newItem;  
    percolateUp(size);  
}
```

❖ What is the time complexity?

❖ $O(\log n)$

Aside: Decrease Key

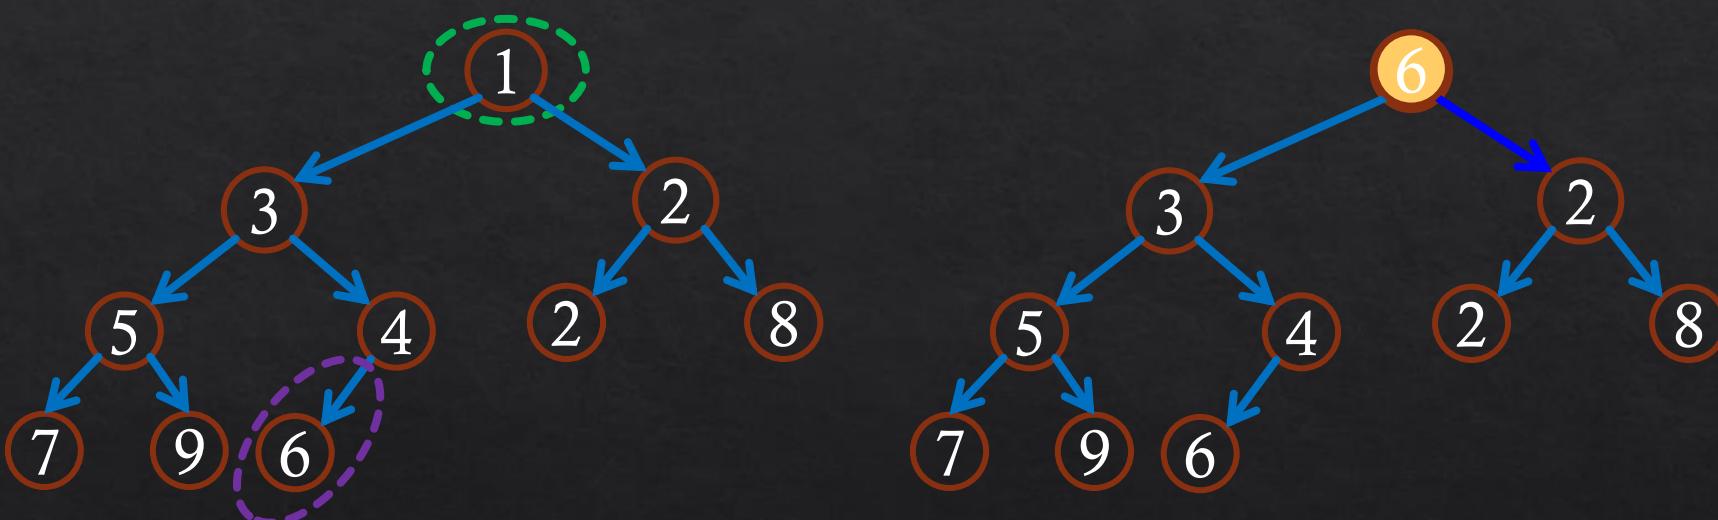
- ❖ Percolating-up can also be exploited to implement the decreasing-key operation



Procedure of dequeueMin

- ❖ The min item is at the root. Save that item to be returned.
- ❖ Move the item in the rightmost leaf of the tree to the root.

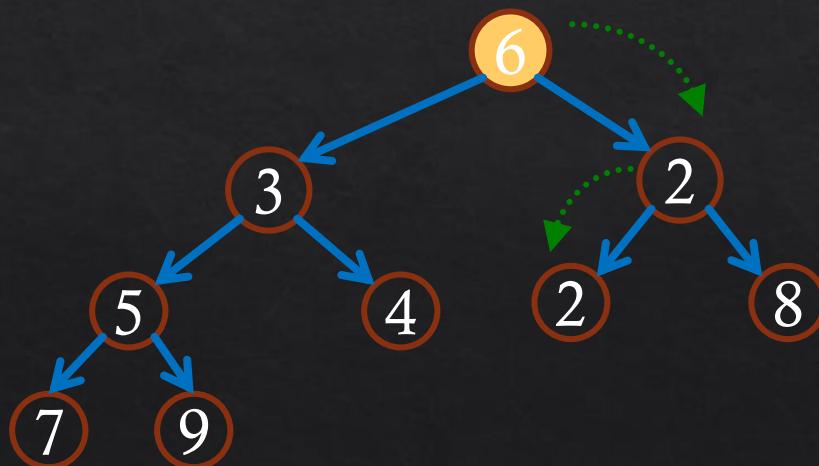
```
swap(heap[1], heap[size--]);
```



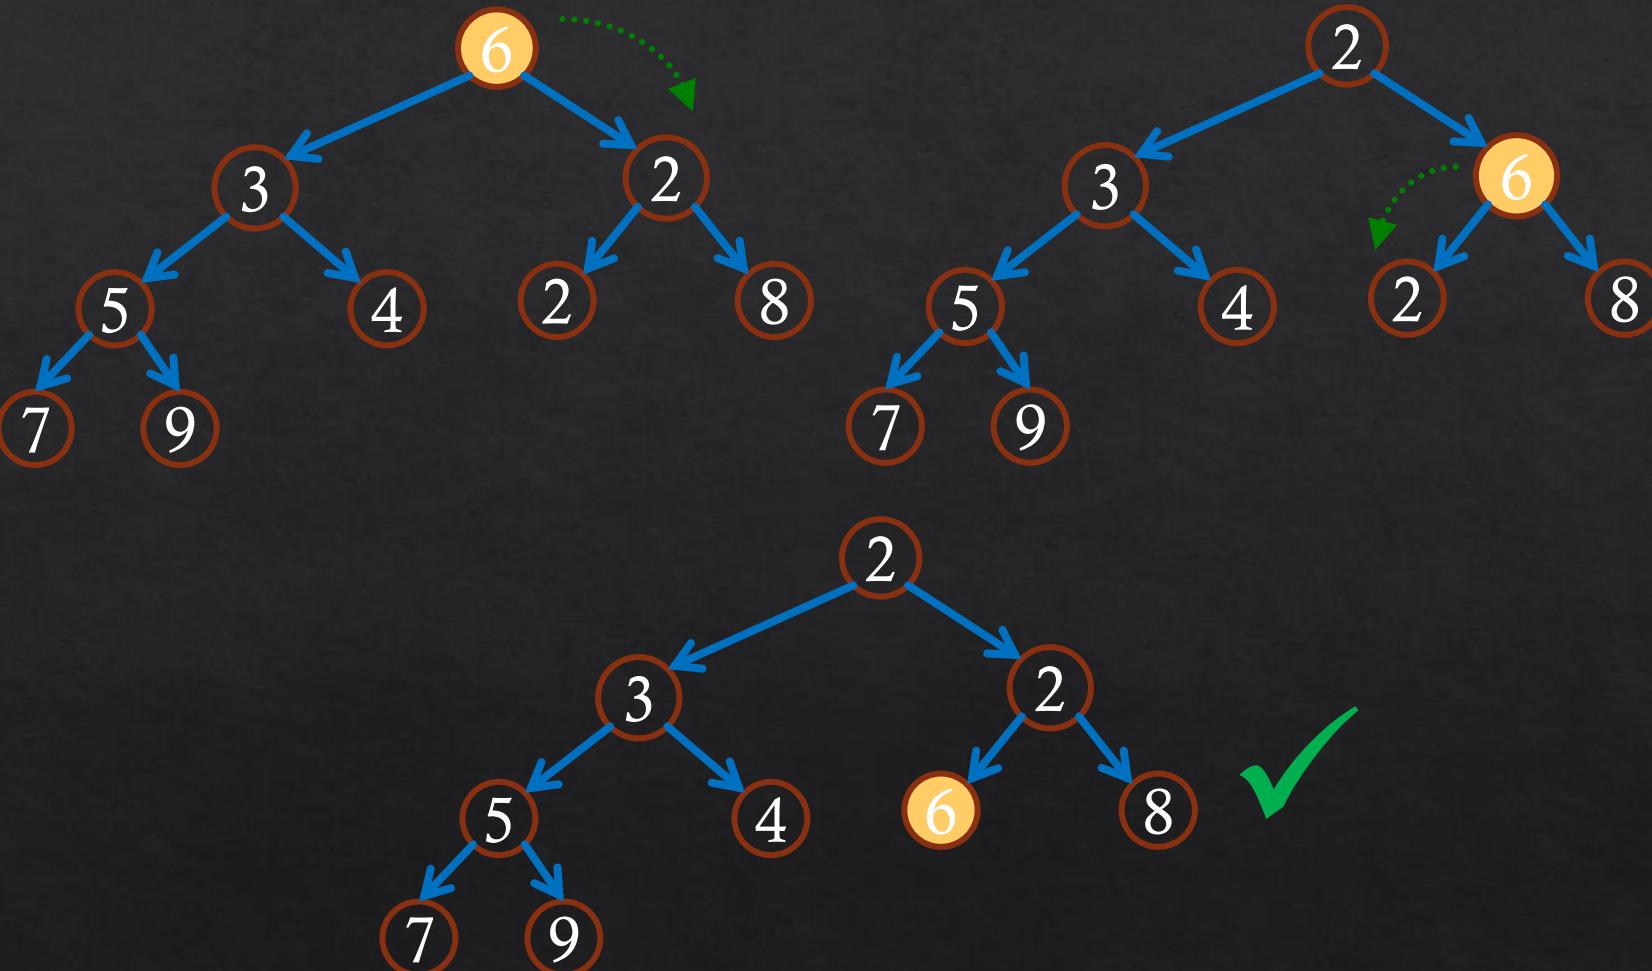
- ❖ The tree may no longer be a heap at this point!

Procedure of dequeueMin

- ❖ **Percolate down** the recently moved item at the root to its proper place to restore heap property.
 - ❖ For each subtree, if the root has a **larger** search key than **either of its children**, swap the item in the root with that of the **smaller** child.



Percolate Down Illustration



Percolate Down Code

```
void minHeap::percolateDown(int id) {  
    for(j = 2*id; j <= size; j = 2*j) {  
        if(j < size && heap[j] > heap[j+1]) j++;  
        if(heap[id] <= heap[j]) break;           find the smaller child  
        swap(heap[id], heap[j]);  
        id = j;  
    }  
}
```

- ❖ Pass index (**id**) of array element that needs to be percolated down.
- ❖ Swap the key in the given node with the smallest key among the node's children, moving down to that child, until:
 - ❖ we reach a leaf node, or
 - ❖ both children have larger (or equal) key

dequeueMin

Code

```
Item minHeap::dequeueMin() {  
    swap(heap[1], heap[size--]);  
    percolateDown(1);  
    return heap[size+1];  
}
```

- ❖ What is the time complexity?
 - ❖ $O(\log n)$

Outline

- ❖ Priority Queue
- ❖ Min Heap and Its Operations
- ❖ Min Heap Initialization and Application

Initializing a Min Heap

- ❖ How do we initialize a min heap from a set of items?
- ❖ Simple solution: insert each entry one by one.
 - ❖ The worst case time complexity for inserting the k -th item is $O(\log k)$, so creating a heap in this way is $O(n \log n)$.
- ❖ Instead, we can do better by putting the entries into a **complete** binary tree and running **percolate down** intelligently.

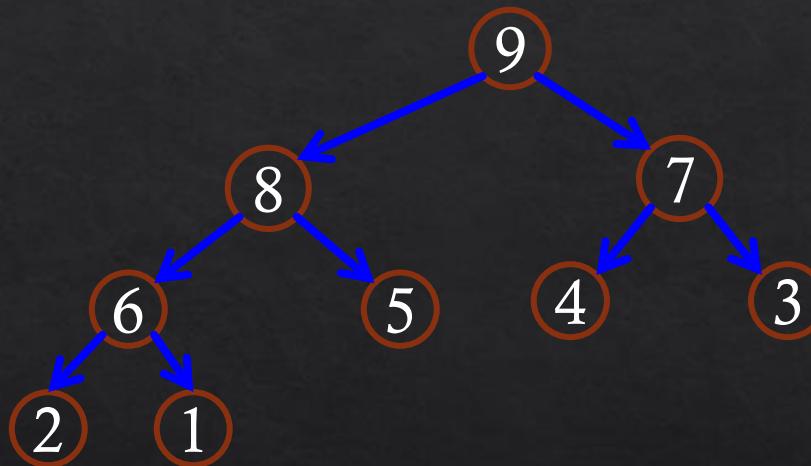
Initializing a Min Heap

- ❖ Put all the items into a complete binary tree.
 - ❖ Implemented using an array.
- ❖ Starting at the rightmost array position **that has a child**, percolate down all nodes in **reverse** level-order.
 - ❖ The rightmost array position **that has a child** is **size/2**.
- ❖ Procedure:

```
For i = size/2 down to 1
percolateDown(i);
```

Initializing a Min Heap Illustration

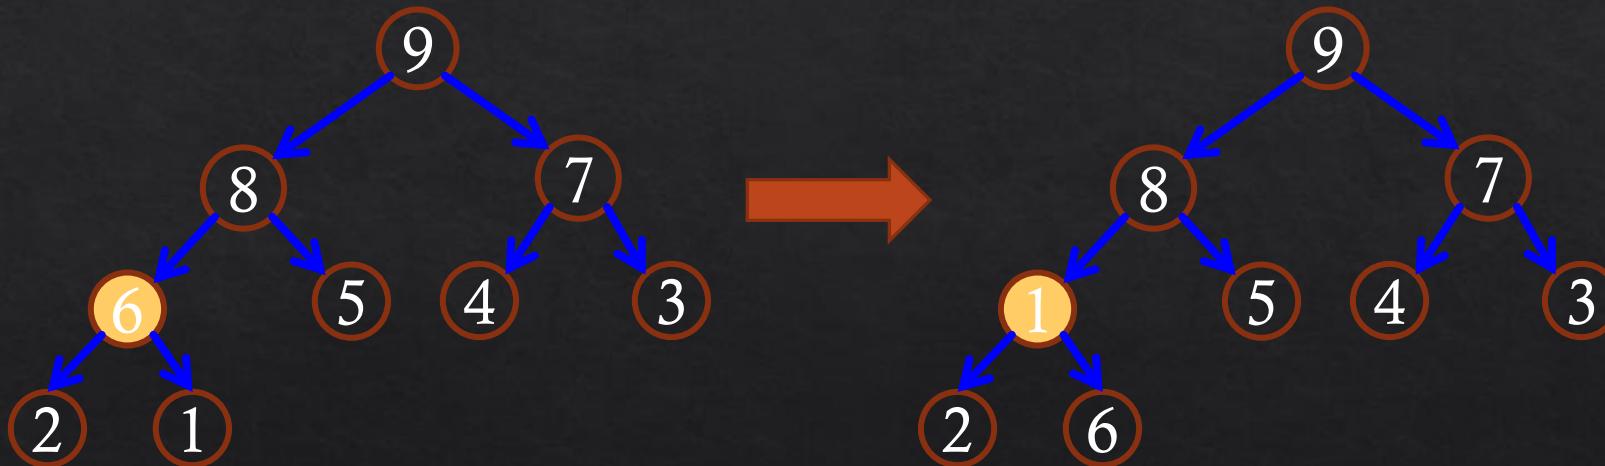
- ❖ Input items: 9, 8, 7, 6, 5, 4, 3, 2, 1
- ❖ First step: put all the items into a complete binary tree.



Initializing a Min Heap Illustration

- ❖ Starting at the rightmost array position **that has a child**, percolate down all nodes in **reverse** level-order.

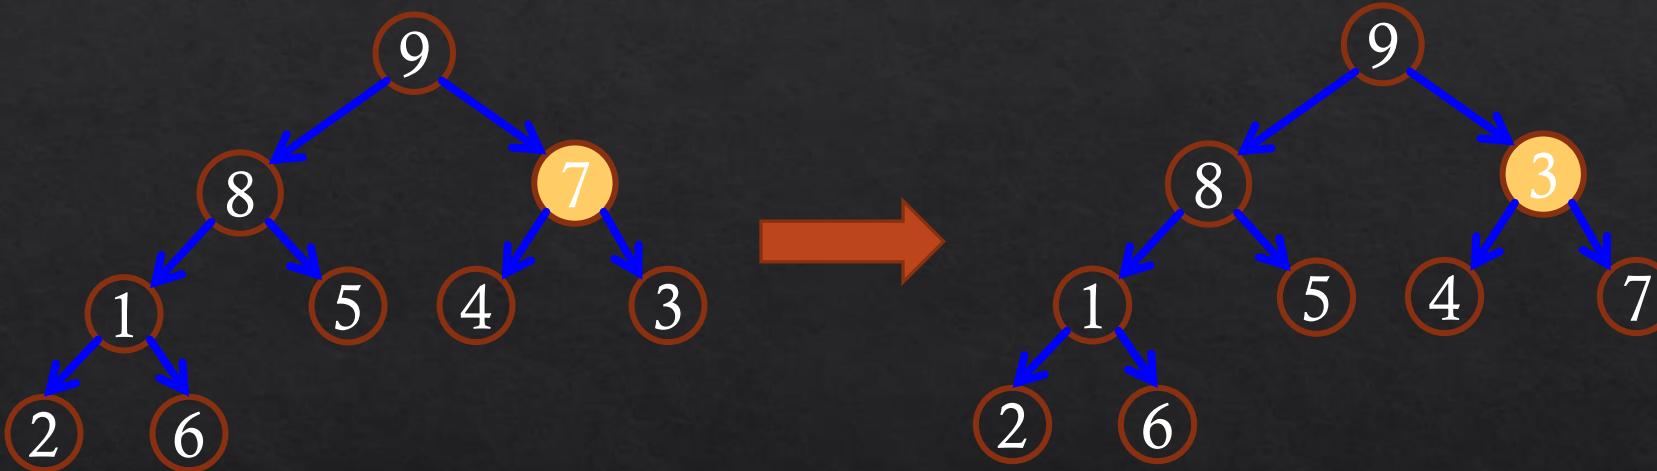
Node at index $9/2 = 4$



Move to next lower array position.

Initializing a Min Heap Illustration

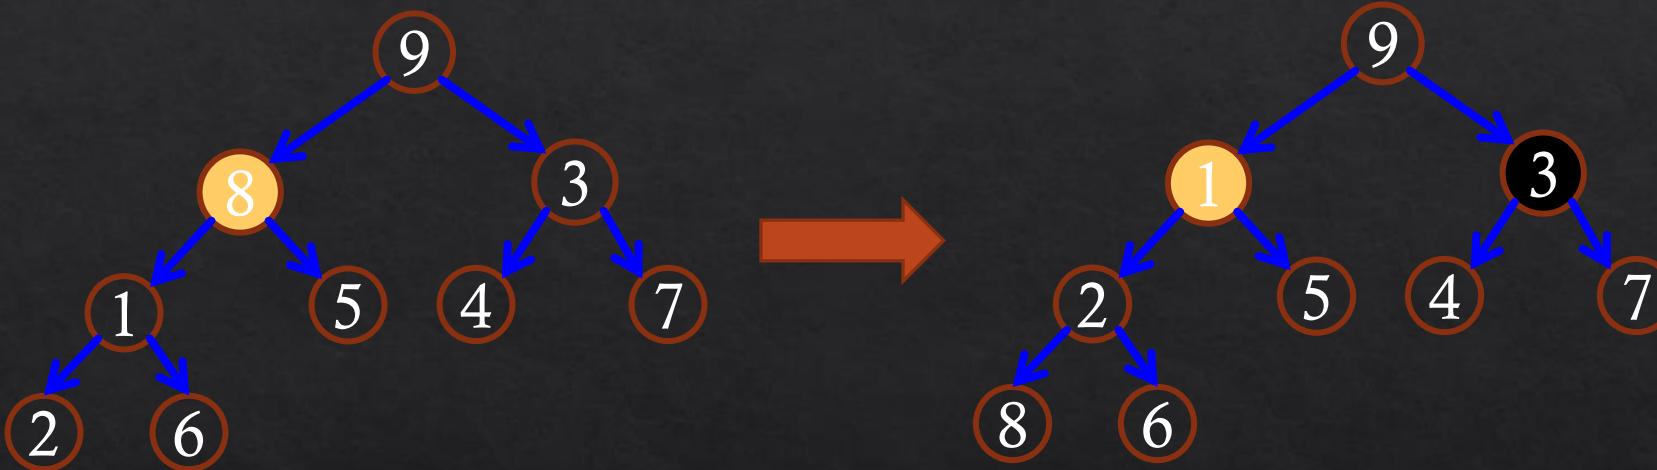
Node at index 3



Move to next lower array position.

Initializing a Min Heap Illustration

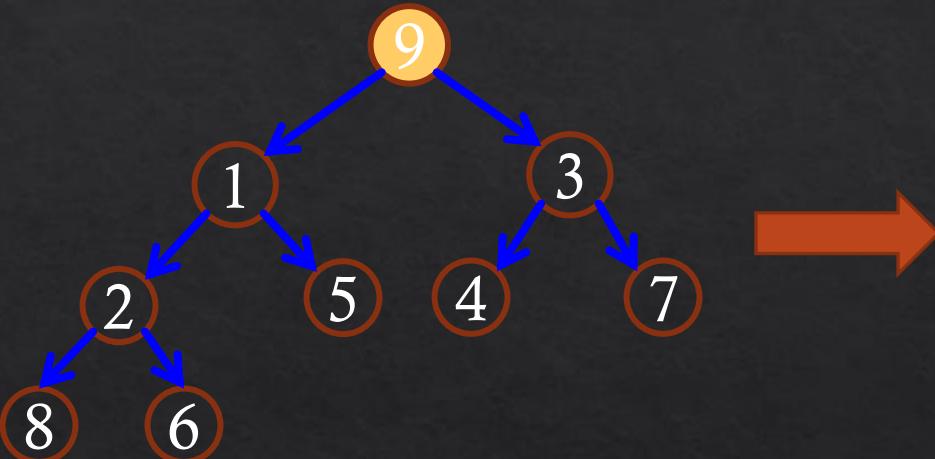
Node at index 2



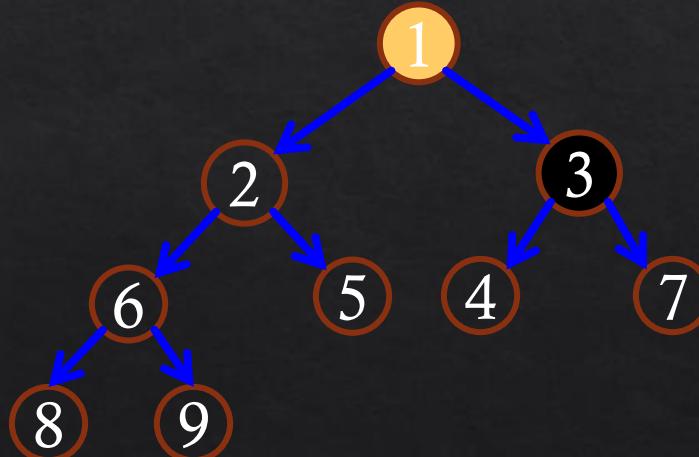
Move to next lower array position.

Initializing a Min Heap Illustration

Node at index 1

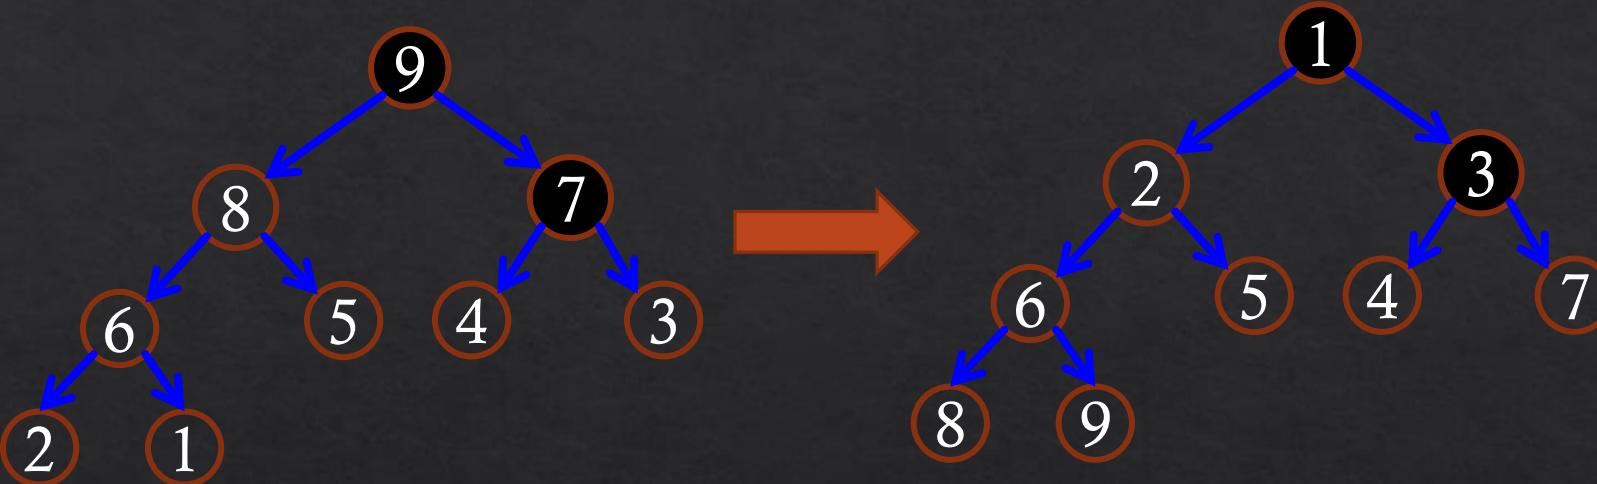


Exercise: What's the result?



Done!

Time Complexity Analysis



- ◊ Suppose: the **height** of the heap is h .
- ◊ Note: Number of nodes at level k ($0 \leq k \leq h$) is $\leq 2^k$.
- ◊ Note: The worst case time complexity of percolating down a node at level k is $O(h - k)$.

Time Complexity Analysis

$$T(h) \leq \sum_{k=0}^{h-1} 2^k O(h - k) = O\left(\sum_{k=0}^{h-1} 2^k (h - k)\right)$$

❖ What is $S(h) = \sum_{k=0}^{h-1} 2^k (h - k)$?

$$S(h) = 2^0 h + 2^1(h - 1) + 2^2(h - 2) + \cdots + 2^{h-1} \cdot 1$$

$$2S(h) = 2^1 h + 2^2(h - 1) + \cdots + 2^{h-1} \cdot 2 + 2^h \cdot 1$$

$$2S(h) = 2^1 h + 2^2(h - 1) + \cdots + 2^{h-1} \cdot 2 + 2^h \cdot 1$$

$$S(h) = 2S(h) - S(h) = 2^1 + 2^2 + \cdots + 2^h - h = 2^{h+1} - 2 - h$$

Time Complexity Analysis

$$T(h) \leq O(2^{h+1} - 2 - h)$$

- ❖ For a complete binary tree, we have

$$h = \lfloor \log_2 n \rfloor \leq \log_2 n$$

where n is the number of nodes.

- ❖ Therefore, the algorithm for initializing a min heap with n nodes has worst case time complexity $T(n) = O(n)$.
 - ❖ Better than the way to enqueue entry one by one.

Application of Heap: Sorting

❖ Procedure:

1. Initialize a min heap with all the elements to be sorted

Complexity: $O(n)$

2. Repeatedly call **dequeueMin** to extract elements out of the heap.

Complexity: $O(n \log n)$

❖ The resulting elements are sorted by their keys.

❖ What is the time complexity?

$O(n \log n)$

❖ This is known as **heap sort**.

Application: Median Maintenance

- ❖ Input: a sequence of numbers x_1, x_2, \dots, x_n , one-by-one
- ❖ Output: at each time step i , the median of x_1, x_2, \dots, x_i
- ❖ Problem: how to do this with $O(\log i)$ time at each step i ?
- ❖ Hint: using two heaps, one min heap and one max heap
- ❖ Key idea: maintain the smallest half ($\left\lfloor \frac{n}{2} \right\rfloor$) in max heap and the largest half ($\left\lceil \frac{n}{2} \right\rceil$) in the min heap
- ❖ Question: How do you get the median (i.e., the $\left\lfloor \frac{n}{2} \right\rfloor$ -th smallest item)?
 - ❖ Answer: get max from the max heap

How to Insert a New Item?

- ❖ Key problem: maintain the **invariant** that the smallest half ($\left\lfloor \frac{n}{2} \right\rfloor$) in max heap and the largest half ($\left\lfloor \frac{n}{2} \right\rfloor$) in the min heap
 - ❖ To maintain balance between the two heaps
 - ❖ If n (before insertion) is even
 - ❖ If new item $\leq \min(\text{minHeap})$, insert it into maxHeap
 - ❖ Else (new item $> \min(\text{minHeap})$), first extract min value from minHeap, then insert that value in maxHeap, and finally insert new item into minHeap
 - ❖ If n (before insertion) is odd
 - ❖ If new item $\geq \max(\text{maxHeap})$, insert it into minHeap
 - ❖ Else (new item $< \max(\text{maxHeap})$), first extract max value from maxHeap, then insert that value in minHeap, and finally insert new item into maxHeap