

VE281

Data Structures and Algorithms

k-d Trees

Learning Objectives:

- Know what a k-d tree is and its difference over basic binary search tree
- Know how to implement search, insertion, and removal for a k-d tree

Multidimensional Search

◊ Example applications:

- ◊ find person by **last name** and **first name** (2D)
- ◊ find location by **latitude** and **longitude** (2D)
- ◊ find book by **author**, **title**, **year published** (3D)
- ◊ find restaurant by **city**, **cuisine**, **popularity**, **price** (4D)

◊ Solution: *k-d tree*

- ◊ $O(\log n)$ insert and search times

1-16 of 143 results for "amd cpu"

Department

- Computer Internal Components
- Computer CPU Processors
- Computer Motherboards
- Internal Fans & Cooling Components
- Computers & Tablets

See All 4 Departments

RESULTS

Customer Reviews

★★★★★ & Up

★★★★★ & Up

★★★★★ & Up

★★★★★ & Up

Brand

- AMD
- Intel

Price

- Under \$25
- \$25 to \$50
- \$50 to \$100
- \$100 to \$200
- \$200 & Above

\$ Min \$ Max Go

CPU Manufacturer

- AMD
- Intel

Computer Processor Type

- AMD A-Series
- AMD Athlon
- AMD FX
- AMD Phenom II
- AMD Ryzen 3

See more

CPU Processor Socket Type

- AMD

Number of CPU Cores

- Single Core
- Dual Core
- Quad Core
- Hexa Core
- Octa Core

CPU Processor Speed

- 1 to 1.59 GHz
- 1.60 to 1.79 GHz
- 1.80 to 1.99 GHz
- 2.00 to 2.49 GHz
- 2.50 to 2.99 GHz
- 3.00 to 3.49 GHz
- 3.50 to 3.99 GHz
- 4.0 GHz & Above

Featured

Price: Low to High

Price: High to Low

Avg. Customer Review

Newest Arrivals

Sponsored

AVGPC Q-Box Series Gaming PC (Q-Box_5700G) 4.6 GHz Max Boost AMD Ryzen 7 5700G 8-Core CPU with Radeon Graphics Cools with 240mm Liquid Cooler 16GB DDR4 3200 500SSD Windows 10 AC WiFi

★★★★★ 255

\$945⁰⁰

Ships to China

Sponsored

AMD Ryzen 5 3600 6-Core, 12-Thread Unlocked Desktop Processor(Tray) with Wraith Stealth Cooler

★★★★★ 3

\$189⁰⁰

Ships to China

Only 13 left in stock - order soon.

Best Seller

AMD Ryzen 5 5600X 6-core, 12-Thread Unlocked Desktop Processor with Wraith Stealth Cooler

★★★★★ 11,996

-19% \$250⁰⁰ \$309.00

Get it Fri, Jul 15 - Mon, Aug 1

Only 1 left in stock - order soon.

More Buying Choices

\$189.49 (92 used & new offers)

Amazon's Choice

AMD Ryzen™ 5 5500 6-Core, 12-Thread Unlocked Desktop Processor with Wraith Stealth Cooler

★★★★★ 82

\$176⁷⁵

Get it Fri, Jul 15 - Mon, Aug 1

Only 15 left in stock - order soon.

More Buying Choices

\$138.98 (37 used & new offers)

AMD Ryzen 9 5900X 12-core, 24-Thread Unlocked Desktop Processor

★★★★★ 4,699

-21% \$450⁰⁰ \$569.99

Ships to China

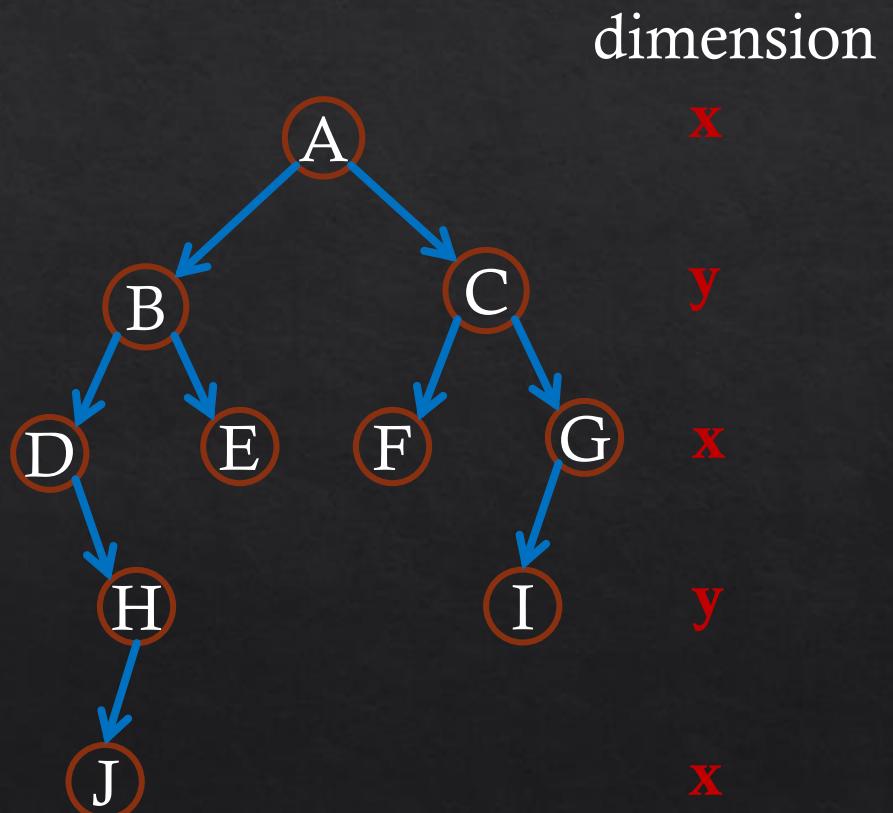
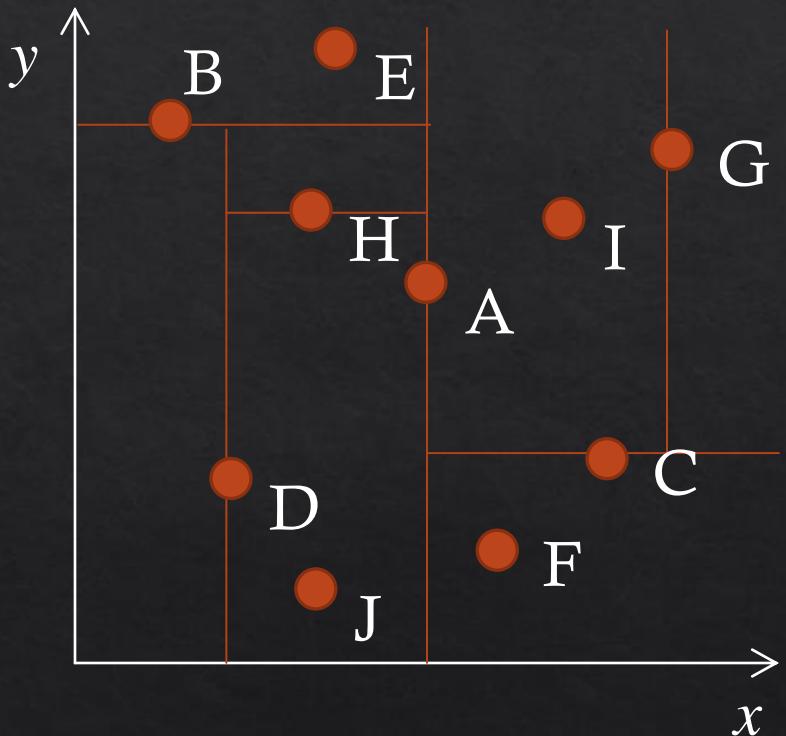
More Buying Choices

k-d Tree

- ◊ A k-d tree is a **binary search tree**
 - ◊ At each level, keys from a different search dimension is used as the **discriminator**
 - ◊ Nodes on the left subtree of a node have keys with value < the node's key value **along this dimension**
 - ◊ Nodes on the right subtree have keys with value \geq the node's key value **along this dimension**
 - ◊ We **cycle** through the dimensions as we go down the tree
-
- ◊ For example, given keys consisting of x- and y-coordinates
 - ◊ level 0 discriminates by the **x-coordinate**
 - ◊ level 1 by the **y-coordinate**
 - ◊ level 2 again by the **x-coordinate**
 - ◊ level 3 again by the **y-coordinate**
 - ◊ etc...

Example

◊ k-d tree for points in a 2-D plane



k-d Tree Insert

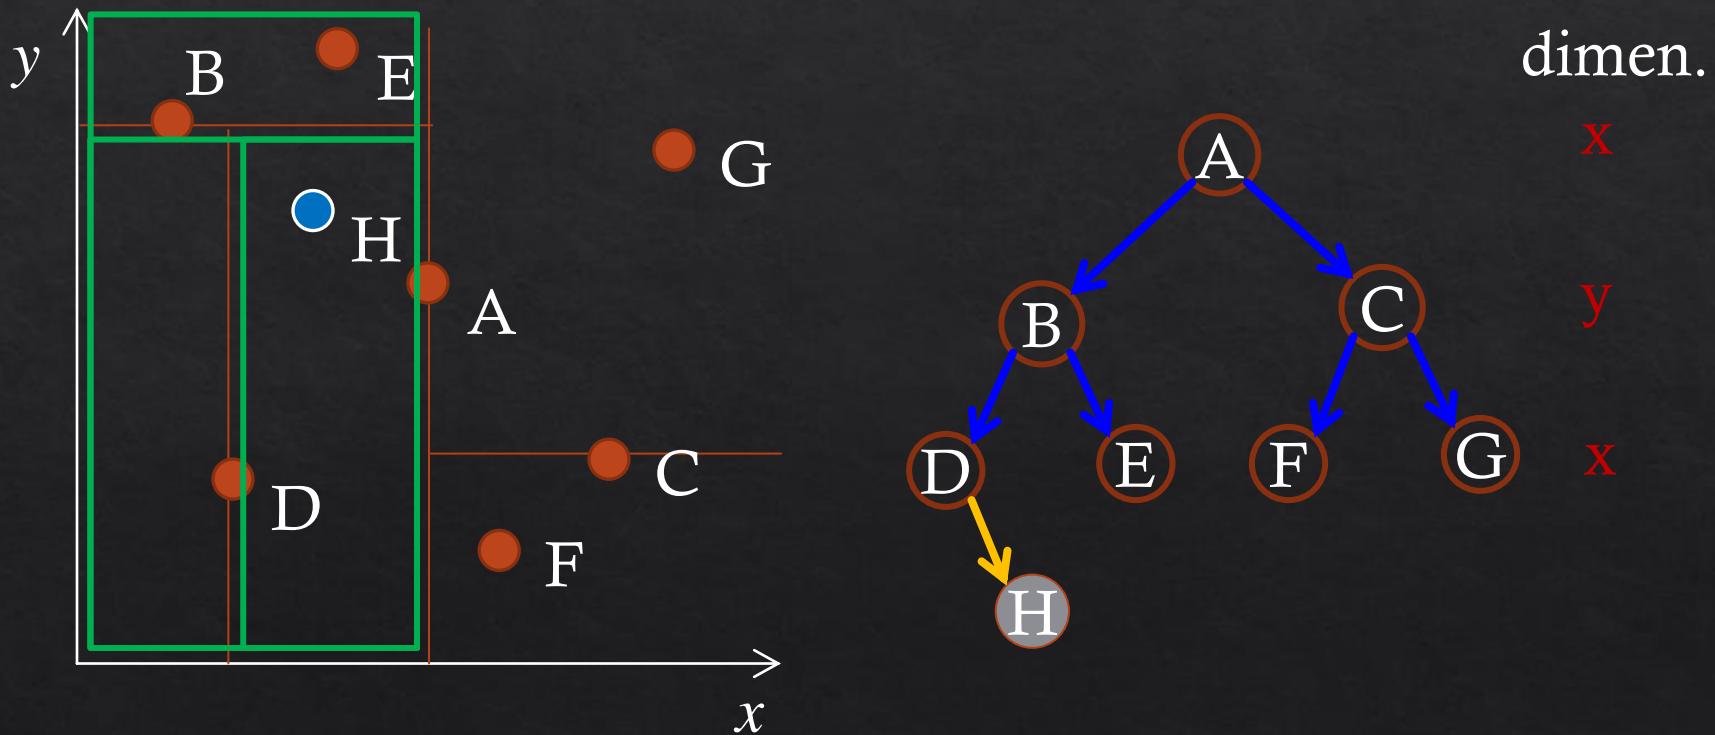
- ◊ If new item's key is equal to the root's key, return;
- ◊ If new item has a key smaller than that of root's along the dimension of the current level, recursive call on left subtree
- ◊ Else, recursive call on the right subtree
- ◊ In recursive call, cyclically increment the dimension

```
void insert(node *&root, Item item, int dim) {  
    if(root == NULL) {  
        root = new node(item);  
        return;  
    }  
    if(item.key == root->item.key) // equal in all  
        return;                                // dimensions  
    if(item.key[dim] < root->item.key[dim])  
        insert(root->left, item, (dim+1)%numDim);  
    else  
        insert(root->right, item, (dim+1)%numDim);  
}
```

dim refers to the dimension of
the root

Insert Example

- ◊ Insert H
- ◊ Initial function call: `insert(A, H, 0)` // 0 indicates dimension x



k-d Tree Search

- ◊ Search works similarly to insert
 - ◊ In recursive call, cyclically increment the dimension

```
node *search(node *root, Key k, int dim) {  
    if(root == NULL) return NULL;  
    if(k == root->item.key)  
        return root;  
    if(k[dim] < root->item.key[dim])  
        return search(root->left, k, (dim+1)%numDim);  
    else  
        return search(root->right, k, (dim+1)%numDim);  
}
```

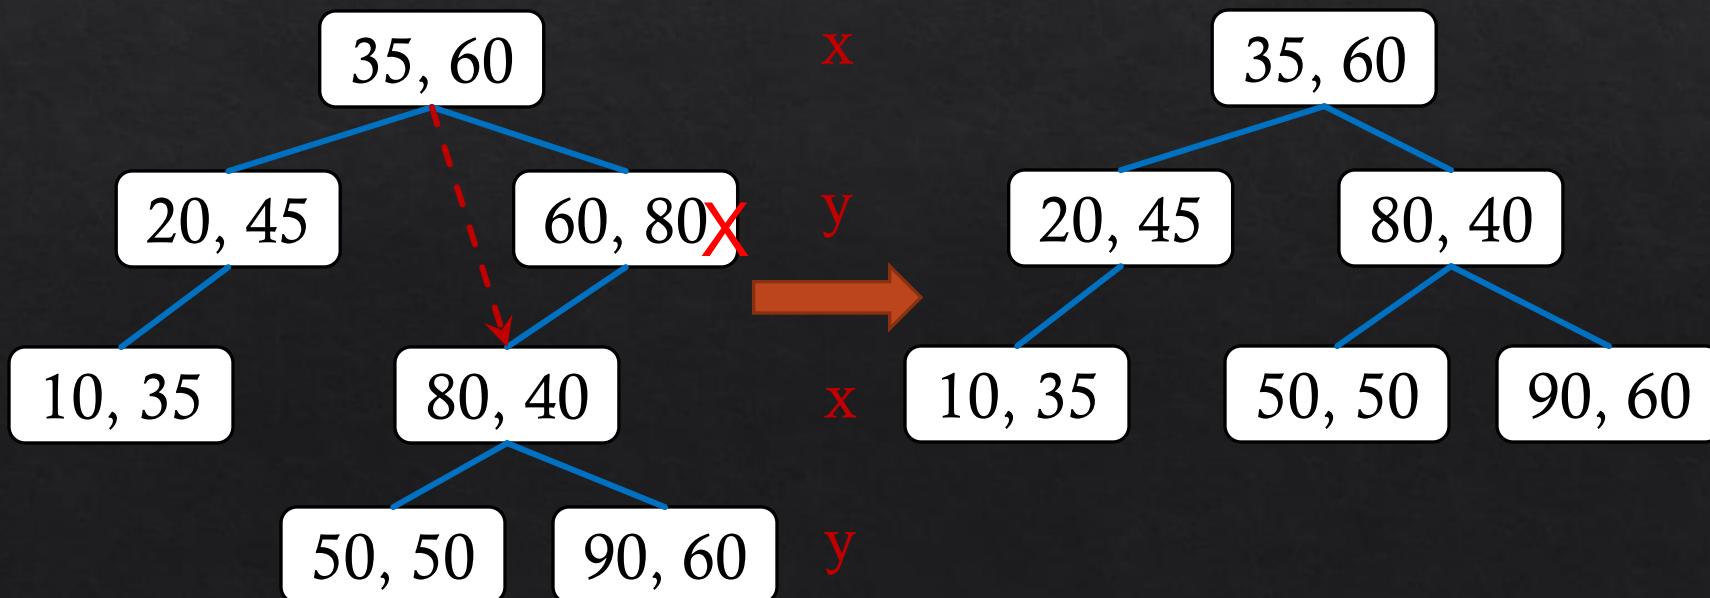
Time complexities of insert and search are all $O(\log n)$

k -d Tree Remove

- ◊ If the node is a leaf, simply remove it (e.g., remove (50,50))
- ◊ If the node has only one child, can we do the same thing as BST (i.e., connect the node's parent to the node's child)?
- ◊ Consider remove (60, 80)

Answer: No!

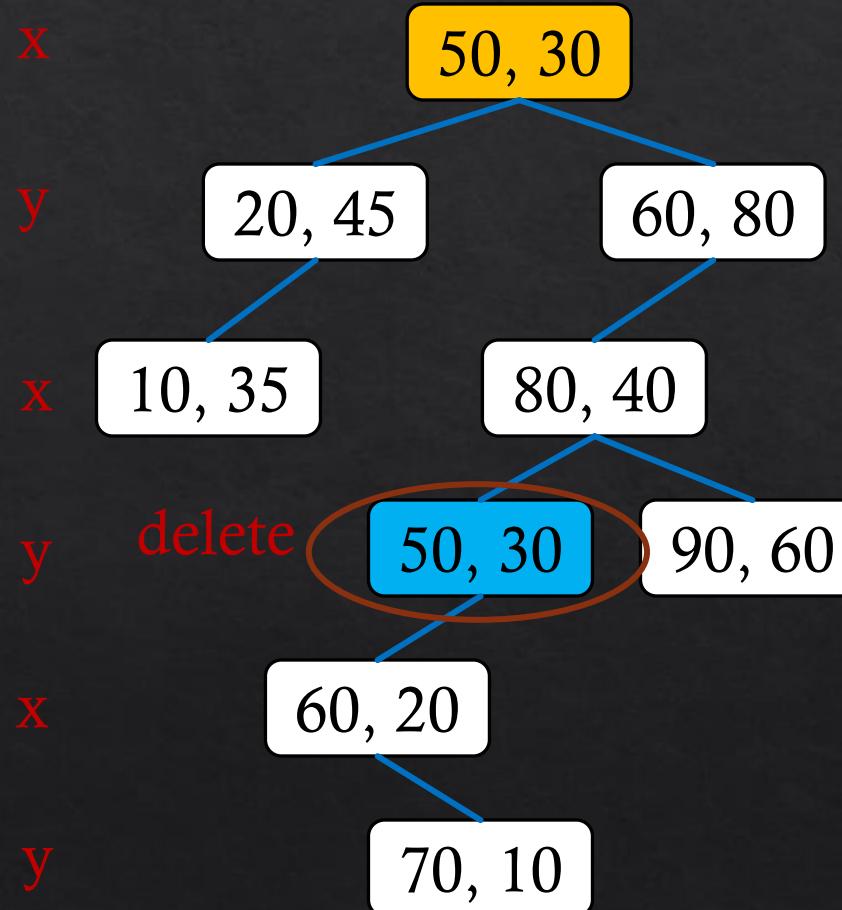
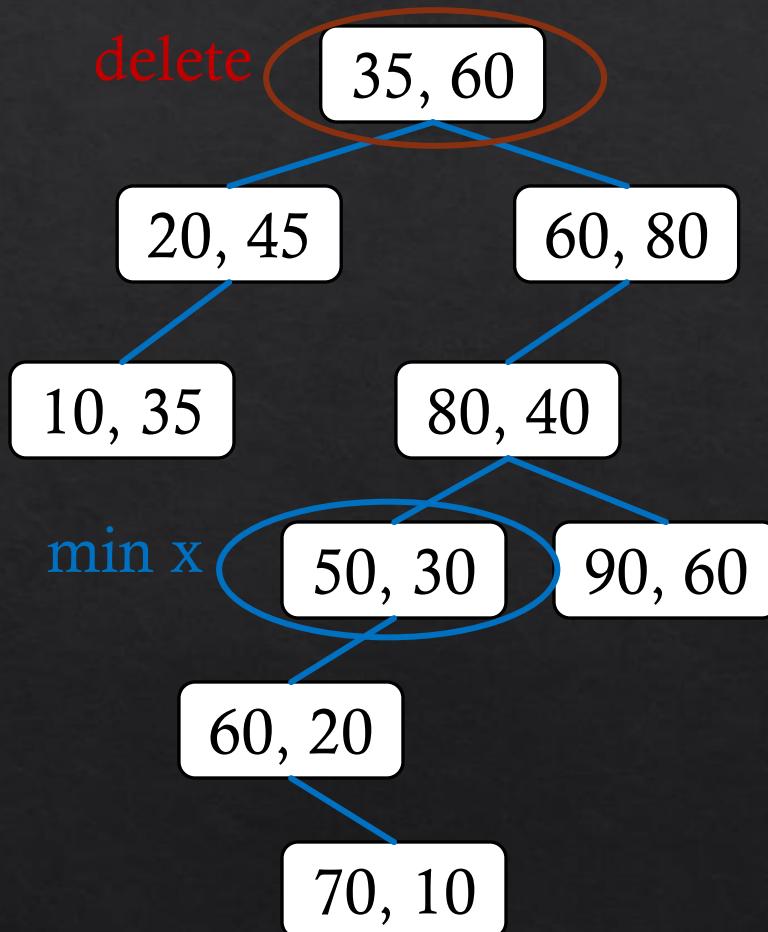
dimension



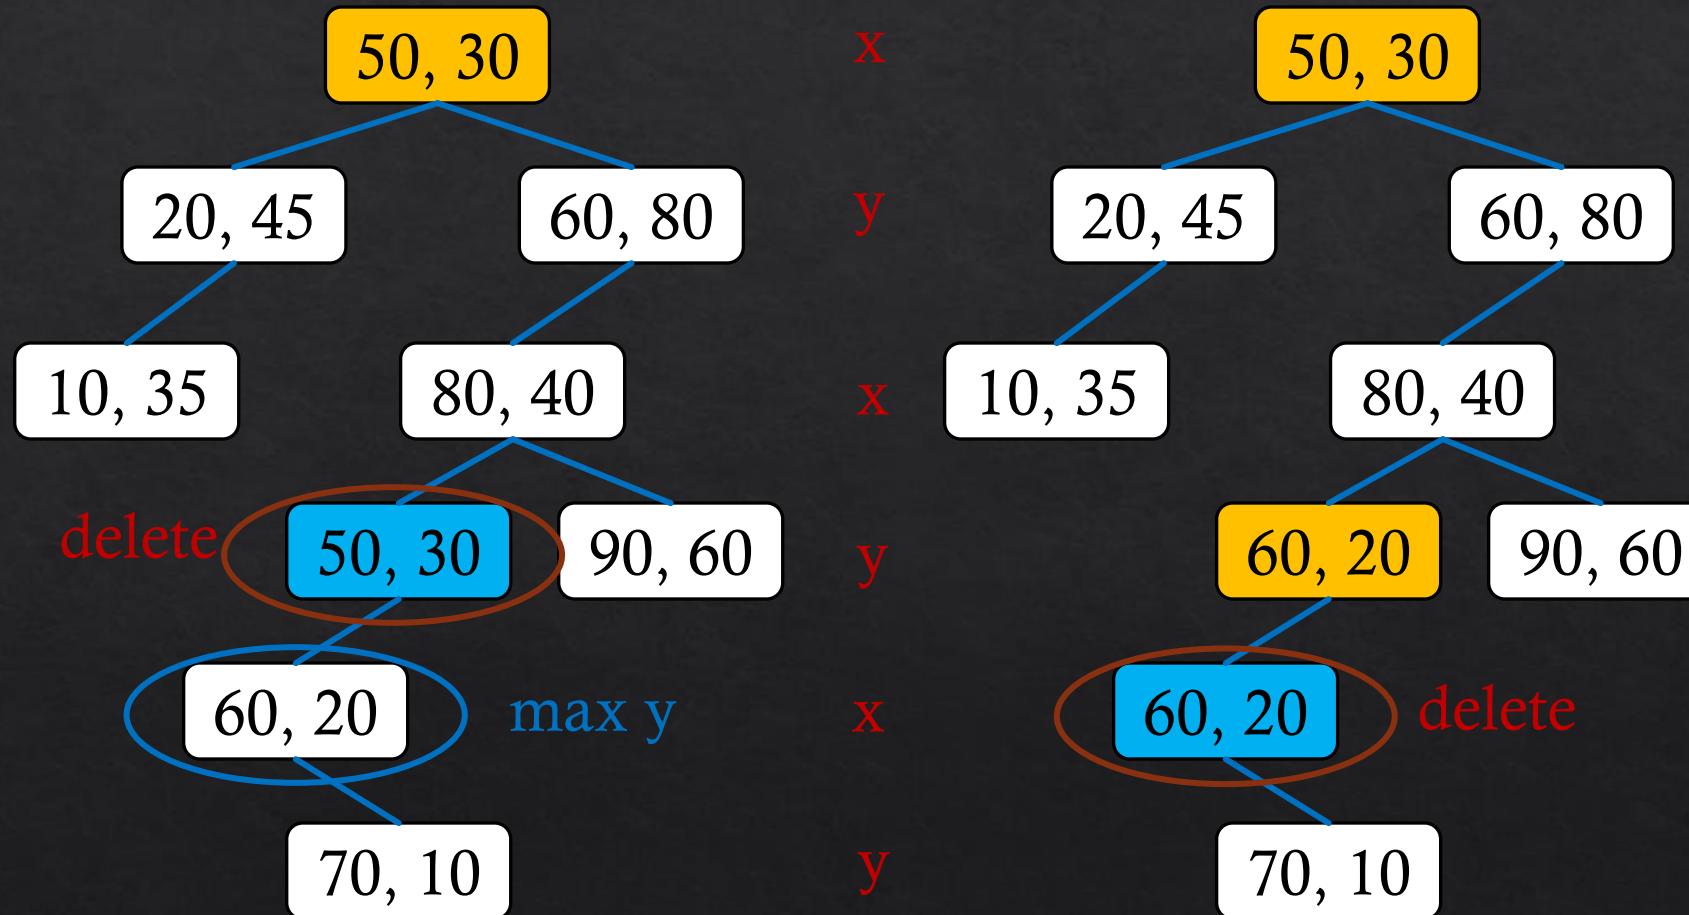
k-d Tree Removal of Non-leaf Node

- ◊ If the node R to be removed has right subtree, find the node M in right subtree with the **minimum** value of the current dimension
 - ◊ Replace the value of R with the value of M
 - ◊ Recurse on M until a leaf is reached. Then remove the leaf
- ◊ Else, find the node M in left subtree with the **maximum** value of the current dimension. Then replace and recurse

k -d Tree Removal Example



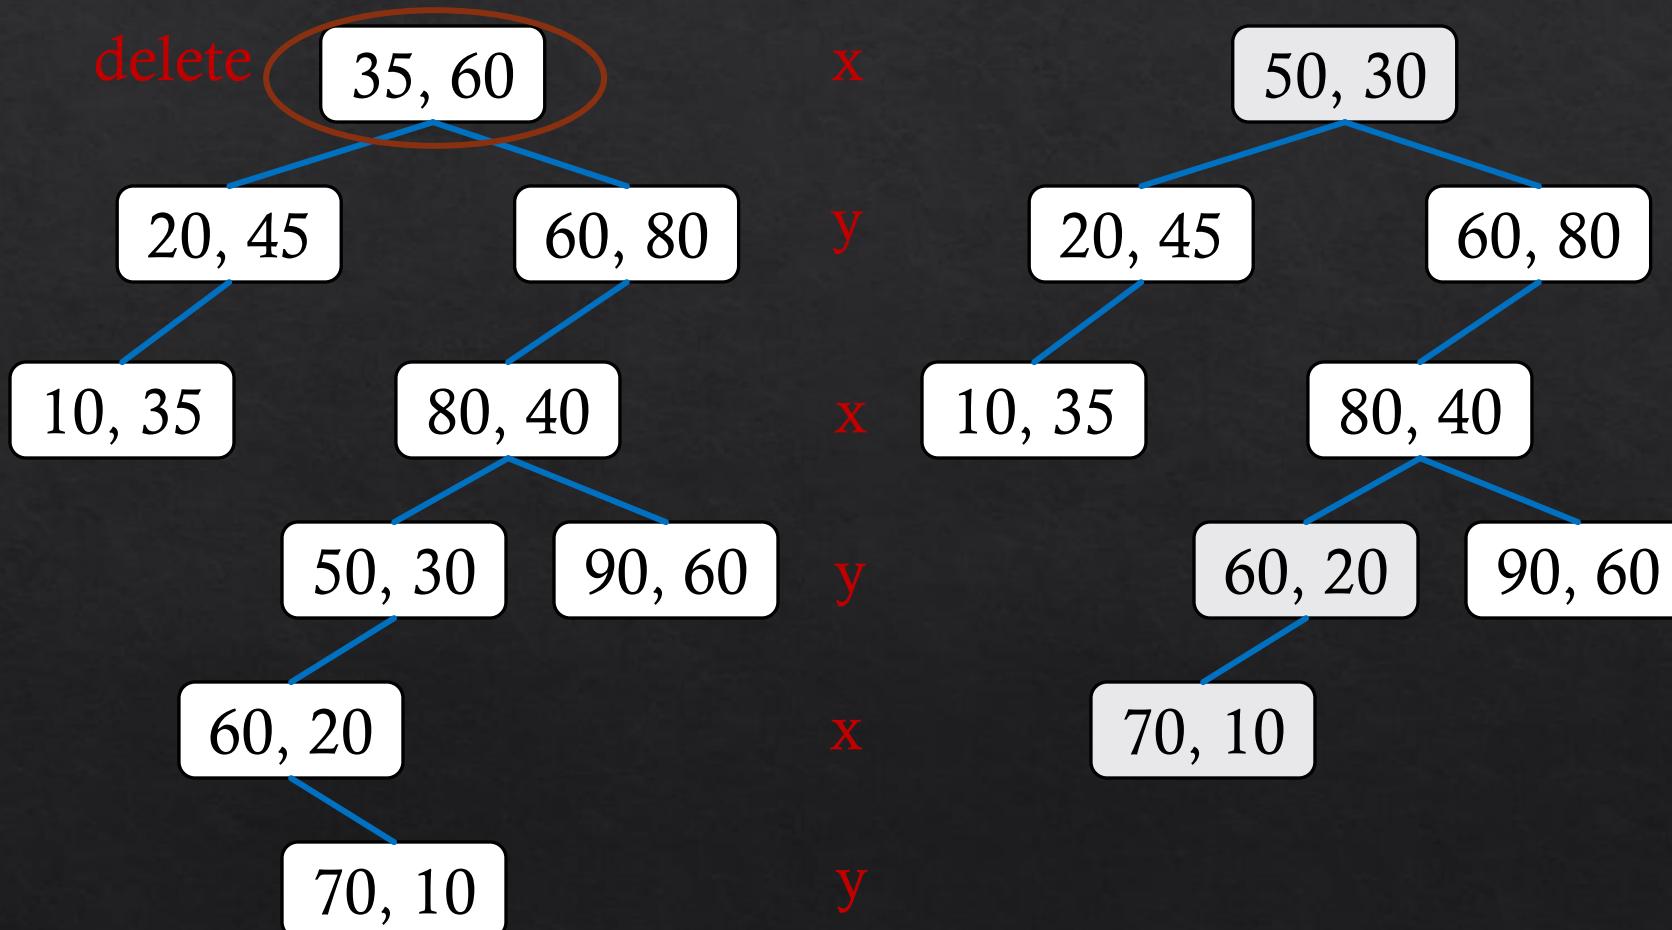
k -d Tree Removal Example



k -d Tree Removal Example

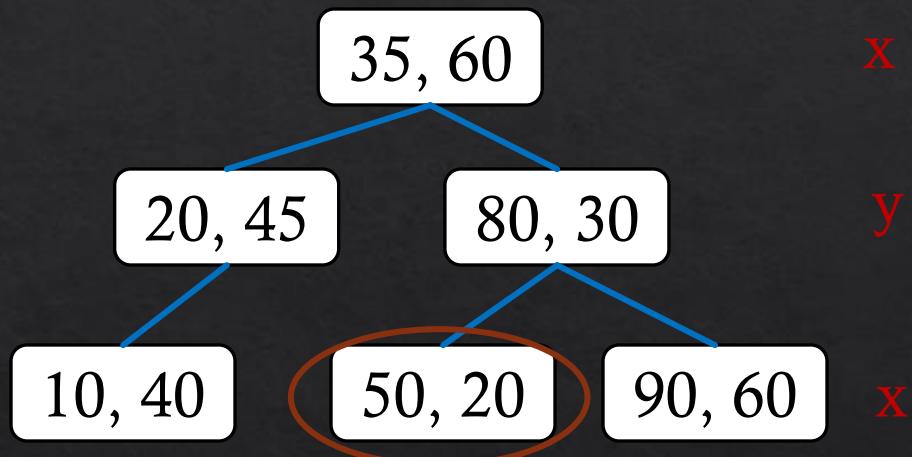


k -d Tree Removal Example: Summary



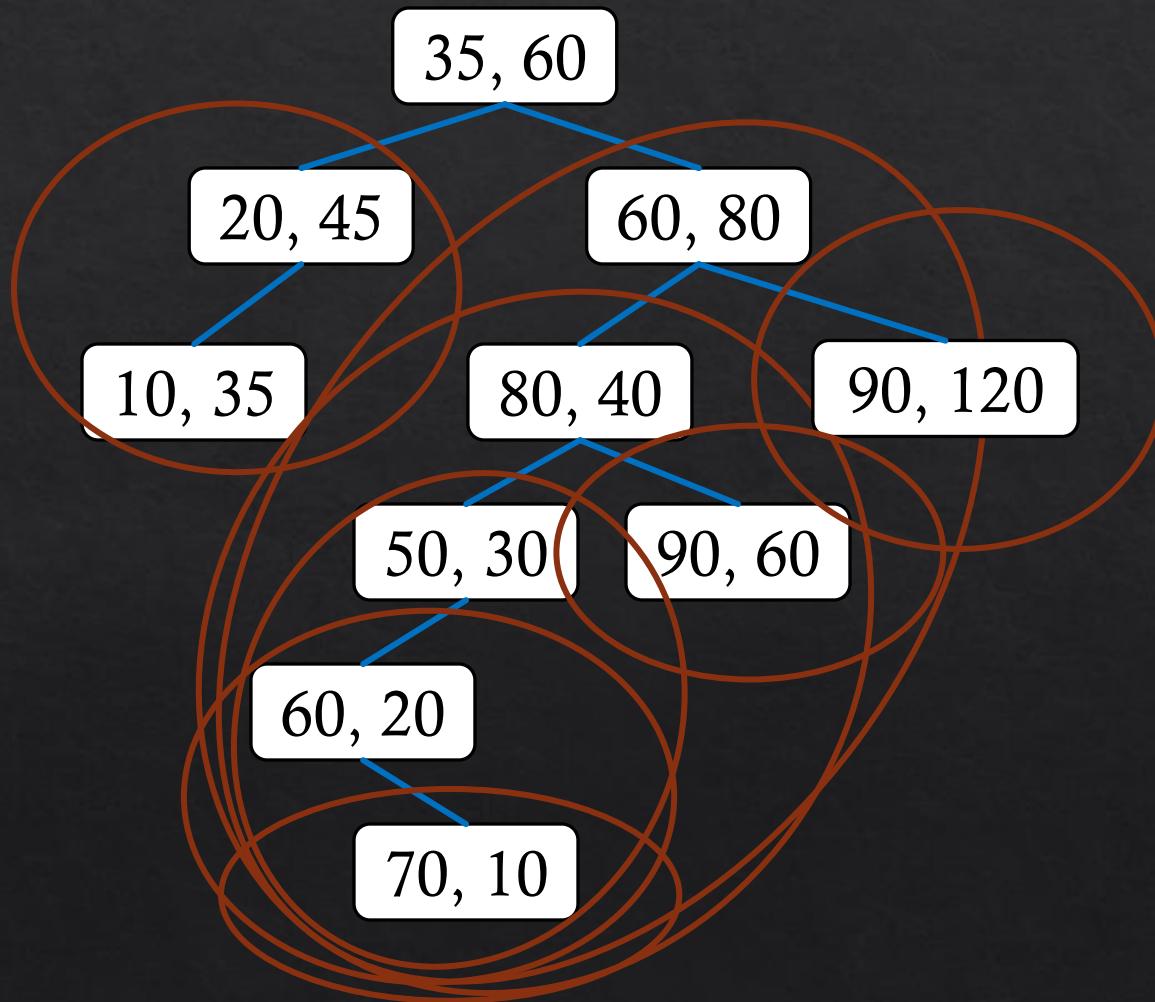
Find Minimum Value in a Dimension

◊ Different from the basic BST, because it may not be the left-most descendent.



Find the node with minimum
value in dimension y

Find Min-Y



Find Minimum Value in a Dimension

```
node *findMin(node *root, int dimCmp, int dim) {  
    // dimCmp: dimension for comparison  
    if(!root) return NULL;  
    node *min = findMin(root->left, dimCmp, (dim+1)%numDim);  
    if(dimCmp != dim) {  
        rightMin = findMin(root->right, dimCmp, (dim+1)%numDim);  
        min = minNode(min, rightMin, dimCmp);  
    }  
    return minNode(min, root, dimCmp);  
}
```

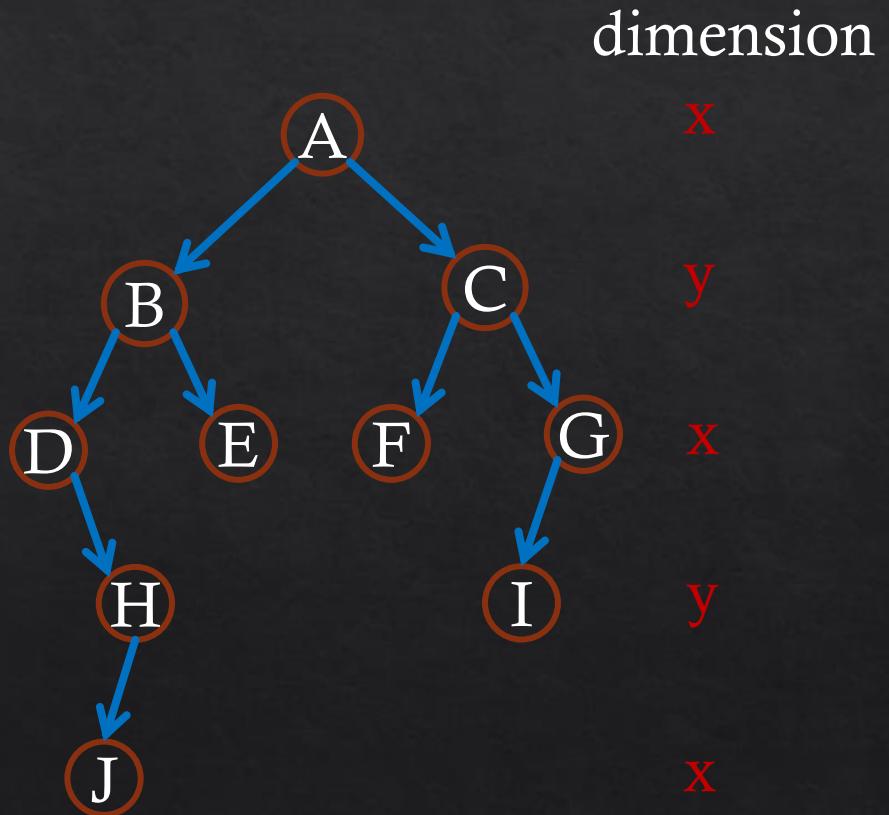
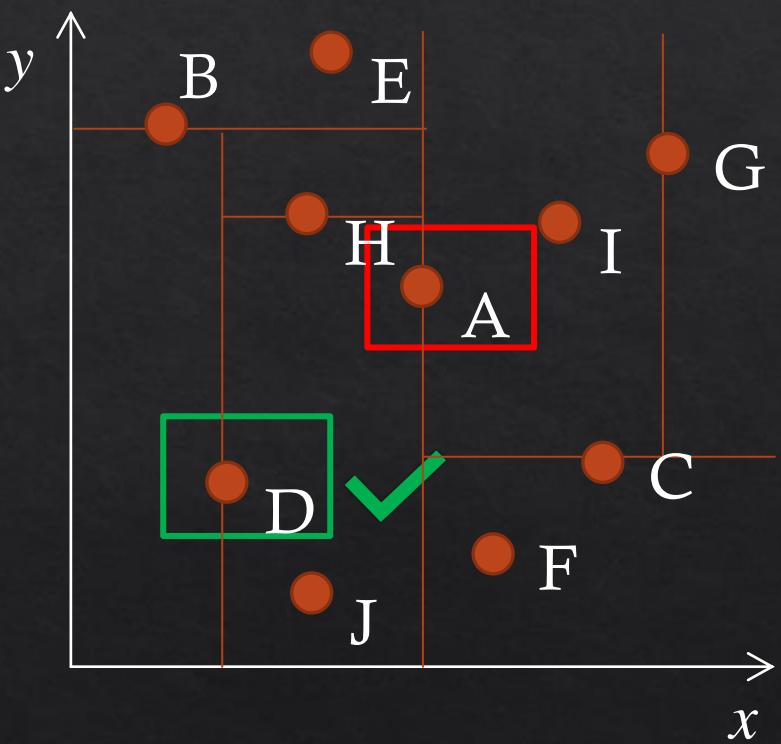
◊**minNode** takes two nodes and a dimension as input, and returns the node with the smaller value in that dimension

Time Complexity of Removal

- ❖ Stop condition of FindMin
 - ❖ A node whose current discriminator is the target dimension
 - ❖ Also the node does not have a left child (no left subtree)

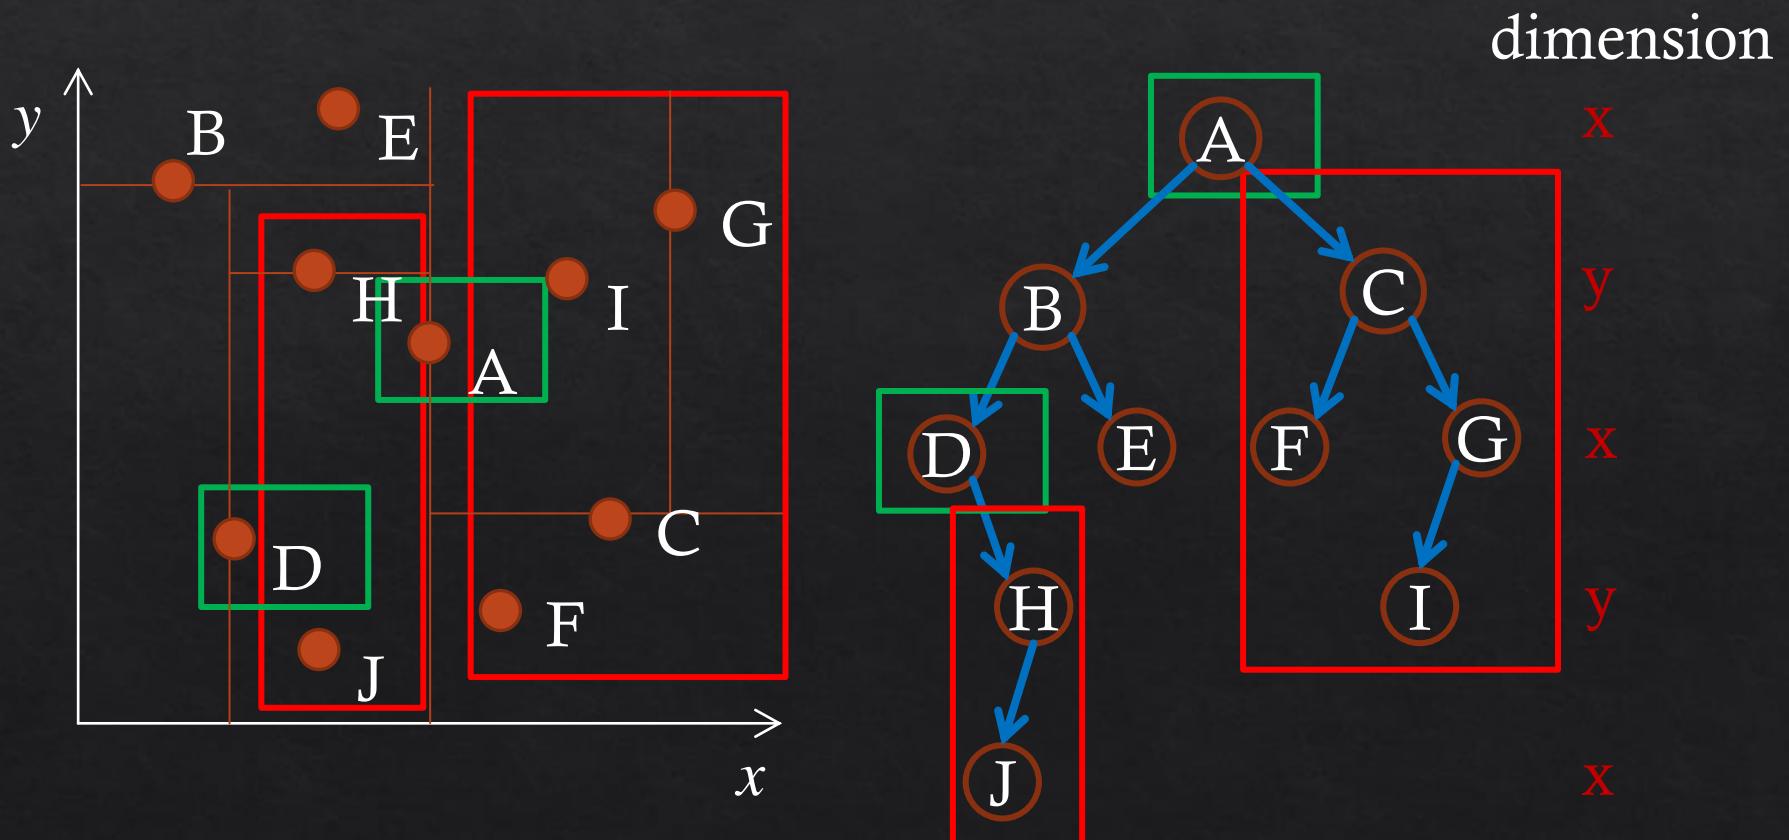
- ❖ Why?
 - ❖ If the node has a left child, the left child < the node

Visual Explanation



Complexity of FindMin

- ◊ FindMin does not explore the right subtree
 - ◊ If the discriminator of the current level is the target dimension
 - ◊ Ignore both the node and the right subtree



A General Analysis

- ❖ If there are M dimensions
- ❖ Nodes are evenly distributed
 - ❖ Prune $\frac{1}{2}$ of the tree in every M levels
- ❖ Assume a total of L levels
- ❖ The whole process touches $(\frac{1}{2})^{L/M}$ Nodes

Cascading FindMin

- ❖ Key idea:
 - ❖ Only cascades to the right subtree
 - ❖ Never searches the same node twice in the same dimension
- ❖ Blackboard!

Multidimensional Range Search

- ❖ Example
 - ❖ Buy ticket for travel between certain dates and certain times
 - ❖ Look for apartments within certain price range, certain districts, and number of bedrooms
 - ❖ Find all restaurants near you

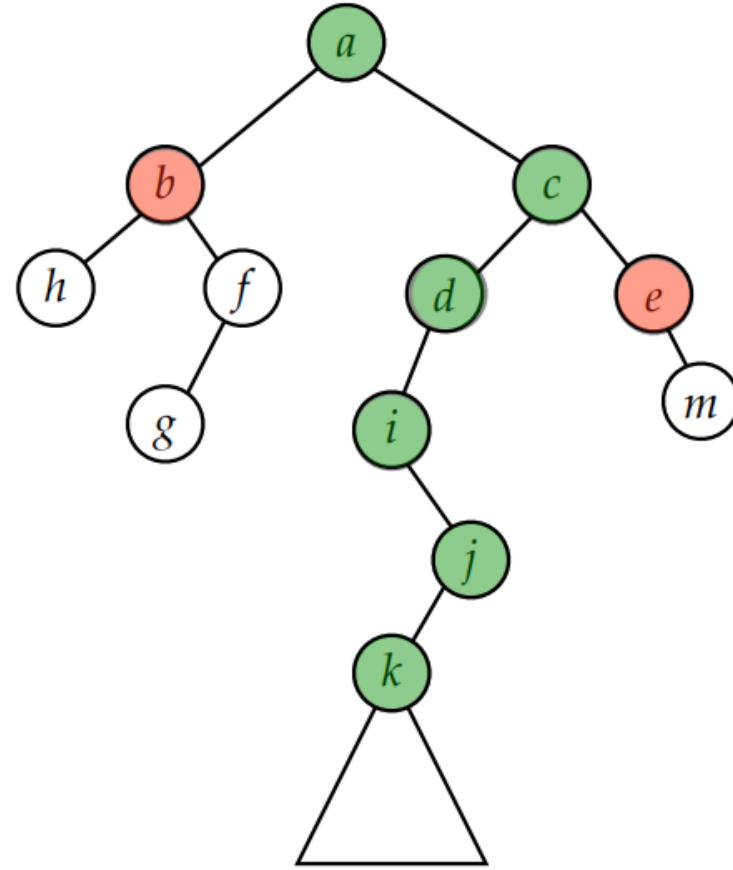
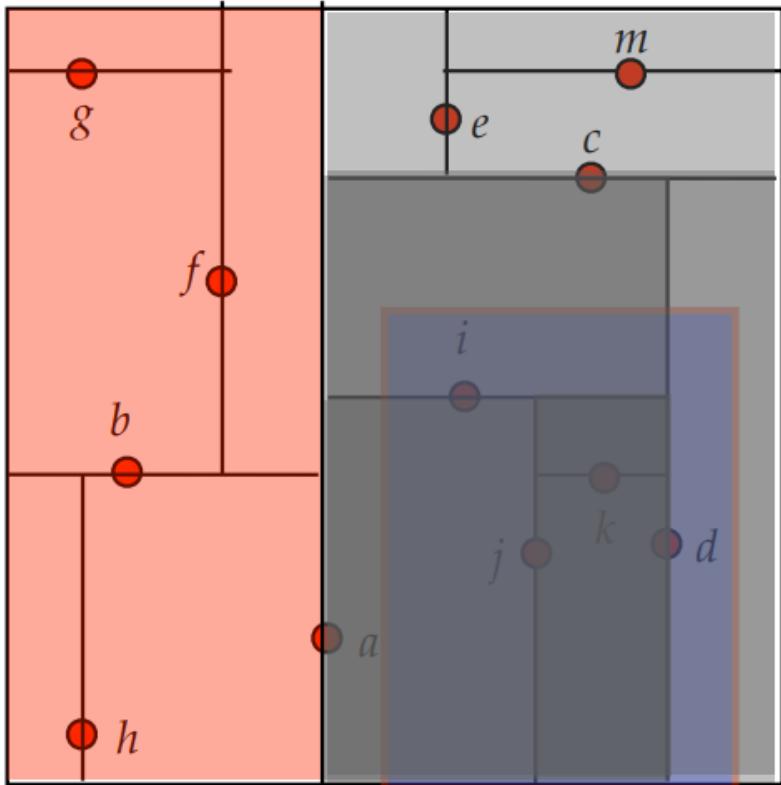
- ❖ k-d tree supports efficient range search, which is similar to that of basic BST but more complex!

k-d Tree Range Search

```
void rangeSearch(node *root, int dim,
    Key searchRange[][][2], Key treeRange[][][2],
    List results)
```

- ❖ Cycle through the dimensions as we go down the level
- ❖ **searchRange**[][][2] holds two values (min, max) per dimension
 - ❖ Define a hyper-cube
 - ❖ min of dimension **j** at **searchRange[j][0]**, max at **searchRange[j][1]**
- ❖ **treeRange**[][][2] holds lower bound and upper bound per dimension for the tree rooted at **root**.
 - ❖ Need to be updated as we go down the levels
 - ❖ Need to check if a search range overlaps a subtree range

Range Searching Example



If query box doesn't overlap bounding box, stop recursion

If bounding box is a subset of query box, report all the points in current subtree

If bounding box overlaps query box, recurse left and right.

Range Query PseudoCode

```
def RangeQuery(Q, T):
    if T == NULL: return empty_set()
    if BB(T) doesn't overlap Query: return 0
    if Query subset of BB(T): return AllNodesUnder(T)

    set = empty_set()
    if T.data in Query: set.union({T.data})

    set.union(RangeQuery(Q, T.left))
    set.union(RangeQuery(Q, T.right))

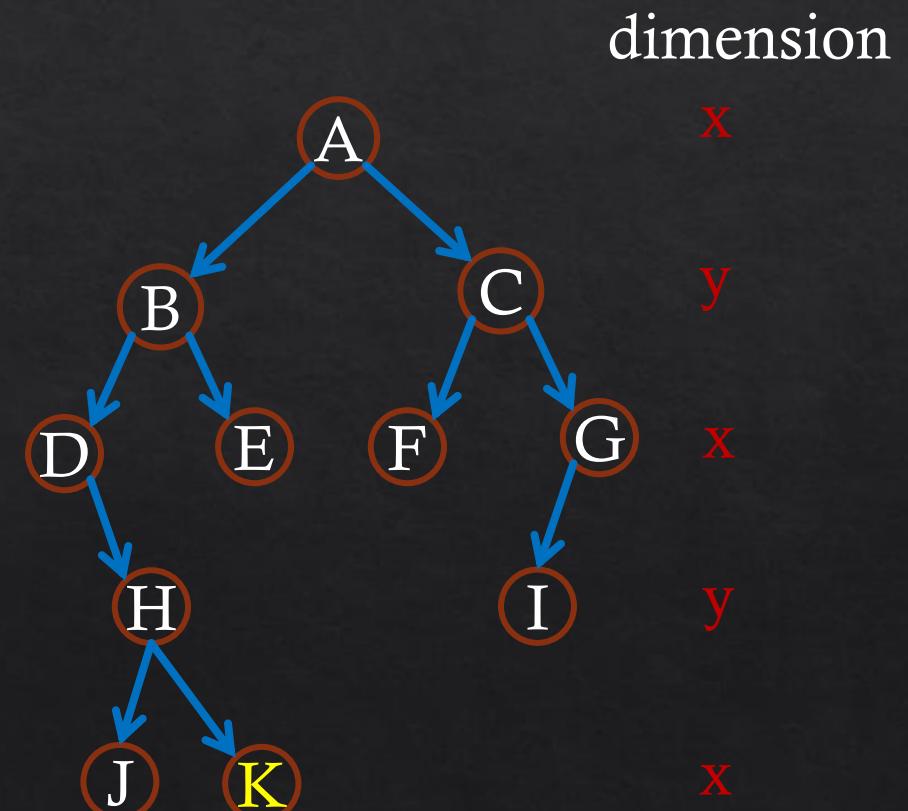
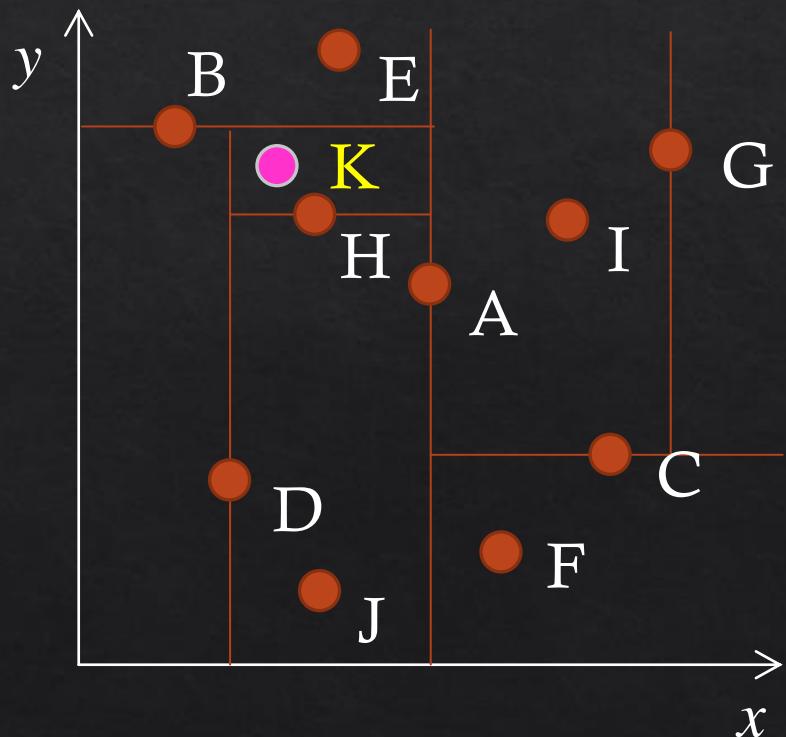
    return set
```

Nearest Neighbor Search

- ◊ Very similar to ranged search.
- ◊ Observation: ranged search is efficient if the range is small.
- ◊ Idea:
 - ◊ Given an element, find a **good** but not **the best** candidate
 - ◊ Outline a small range
 - ◊ Range search in reverse order

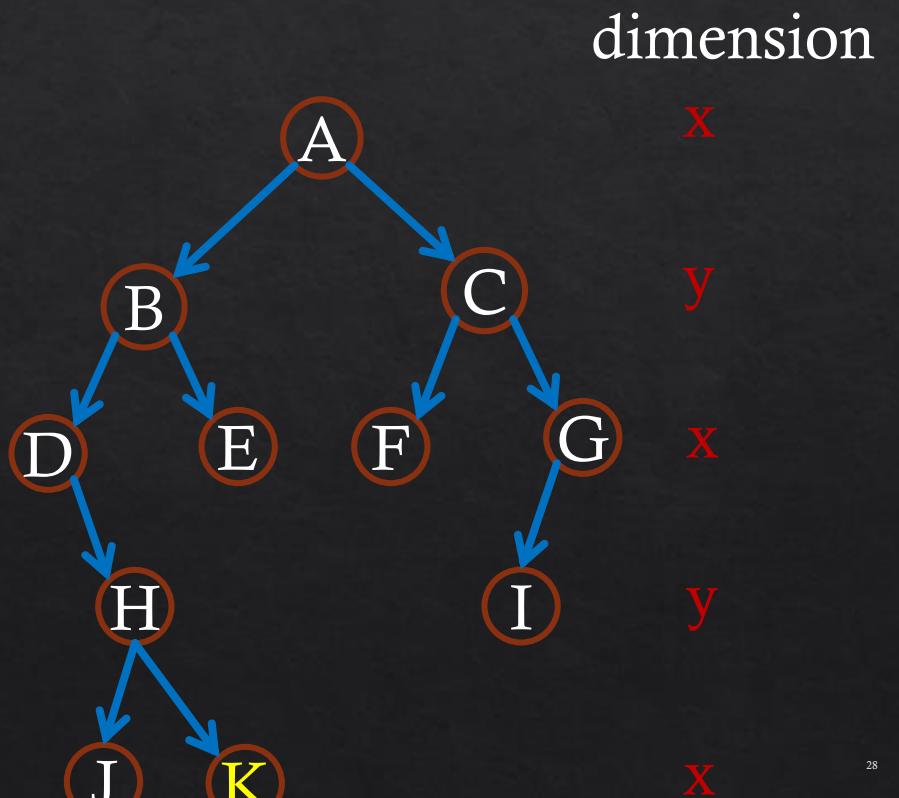
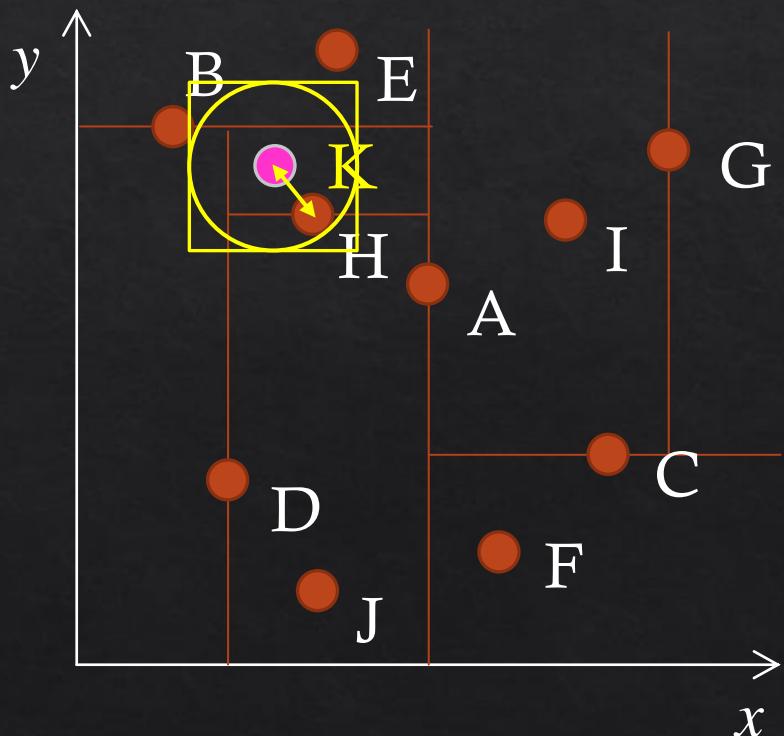
What Is a Good Candidate?

- ◊ Suppose we want to find the closest neighbor of K
- ◊ If we were to add K into the k-d tree
- ◊ Its parent H should be in close vicinity of K
- ◊ H could be a **good** candidate



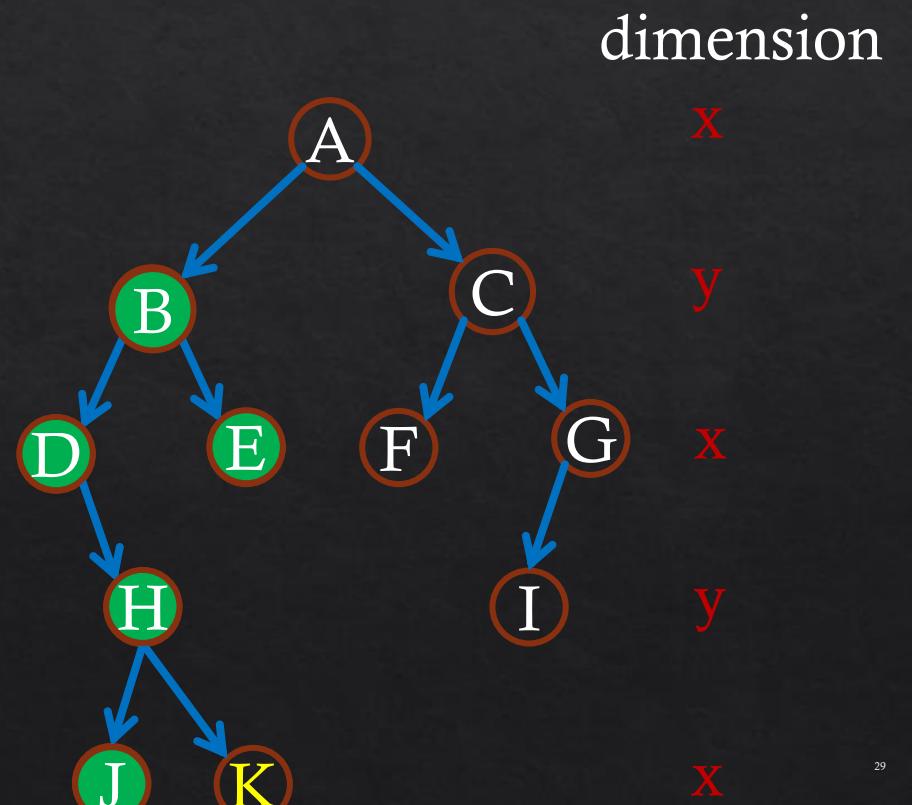
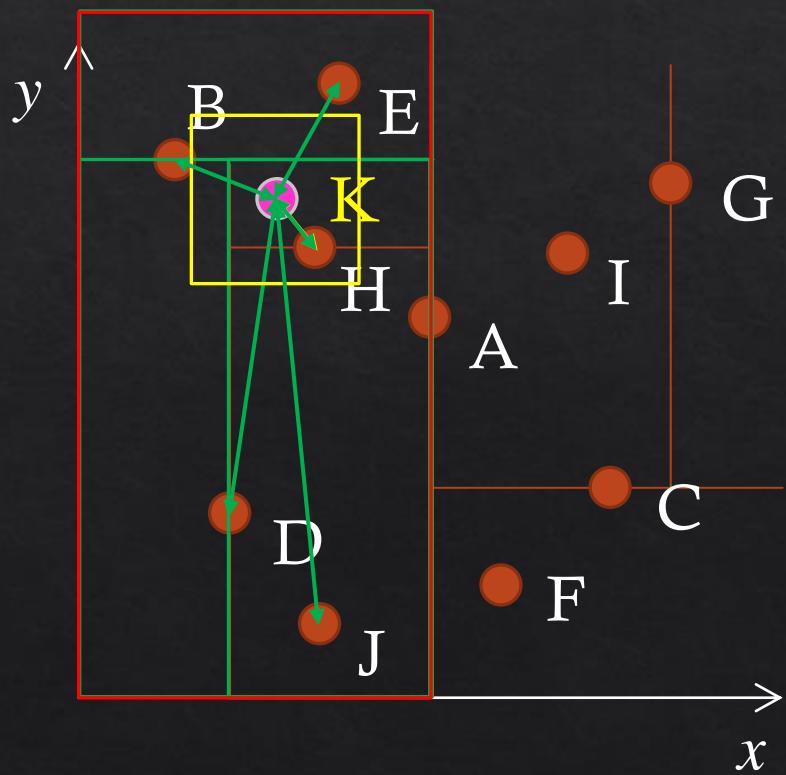
What Is the Range?

- ❖ Compute the Radius
 - ❖ **Better** candidates must locate within the circle (or the sphere)
- ❖ K-d tree can't efficiently search the range of a sphere
- ❖ Set the range as a “rectangle” (or a cuboid in high dimensions)



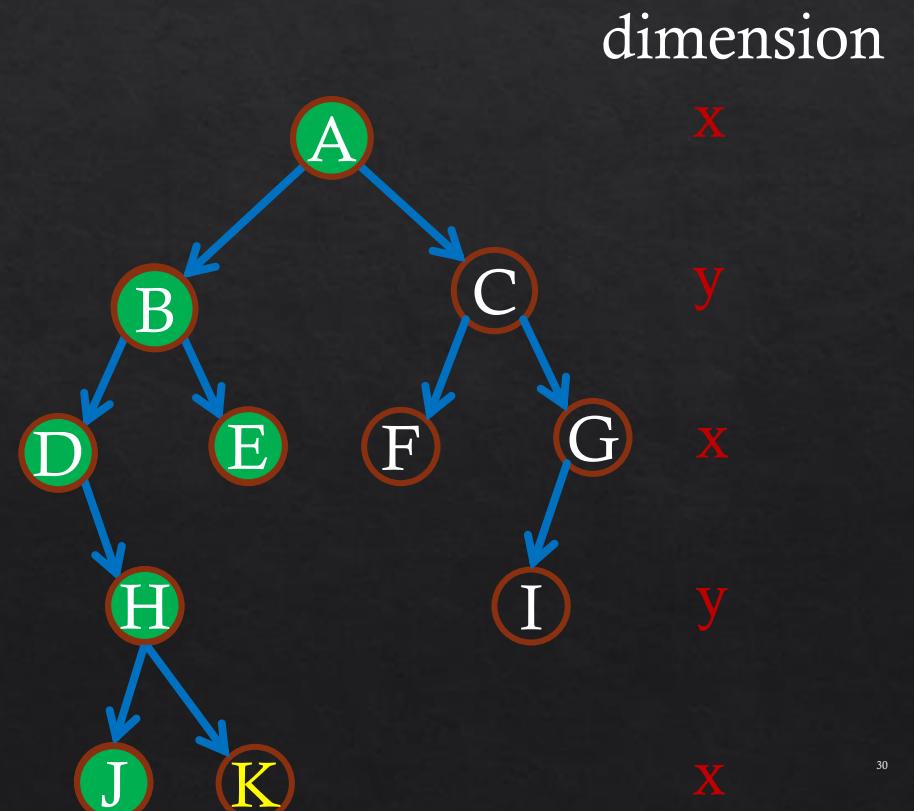
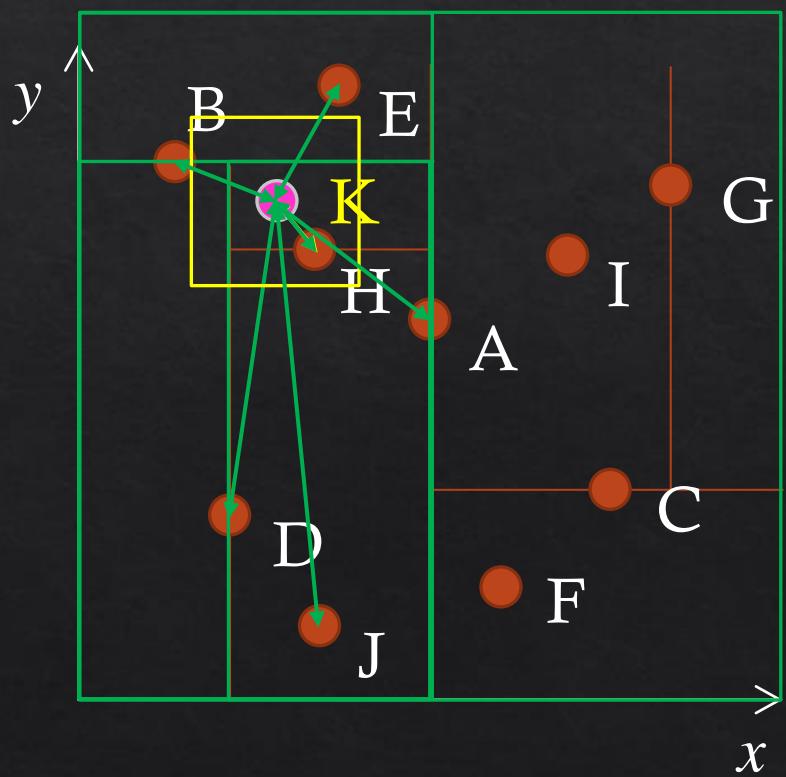
Top down vs Bottom up

- ❖ Bottom up
- ❖ Each node defines a “space”, or a domain
- ❖ Stop when a node completely encompasses the target search space



Bottom up Search

- ◊ Top down
- ◊ Why top down is inefficient:
 - ◊ We start with a small cube already, no need to start big



Implementing Nearest Neighbor Search

- ❖ Modifications to the nodes in k-d tree

```
struct node {  
    node* parent_ptr;  
    vector<int> keys; // Or a key structure  
    Value value;  
    vector<pair<int, int> > domain; // The domain of the current node  
}
```